

Complete Event-Driven Architecture - Beginner’s Guide

A Step-by-Step Explanation of Building Production-Grade Cloud Applications

Author: Learning Guide for Beginners Date: October 2025 Technologies: Apache Kafka, PostgreSQL, MinIO (S3), Java, Docker

Table of Contents

- 1. [Introduction - What Did We Build?](#)
- 2. [The Big Picture - High-Level Architecture](#)
- 3. [Component 1: Apache Kafka - The Message Streaming Platform](#)
- 4. [Component 2: PostgreSQL - The Relational Database](#)
- 5. [Component 3: MinIO \(S3\) - The Blob Storage](#)
- 6. [Component 4: Java Application - The Orchestrator](#)
- 7. [How Everything Works Together](#)
- 8. [Docker & Containerization](#)
- 9. [Deep Dive: Complete Data Flow](#)
- 10. [Real-World Applications](#)
- 11. [Key Concepts for Beginners](#)
- 12. [Troubleshooting & Common Issues](#)

1. Introduction - What Did We Build?

Overview

You’ve just built a **complete event-driven architecture** that’s used by billion-dollar companies like: - **Netflix** - Video streaming and content delivery - **Instagram** - Photo sharing and social media - **Dropbox** - File storage and synchronization - **Spotify** - Music streaming and recommendations - **Uber** - Real-time ride tracking - **Airbnb** - Property listings and bookings

What Makes This “Production-Grade”?

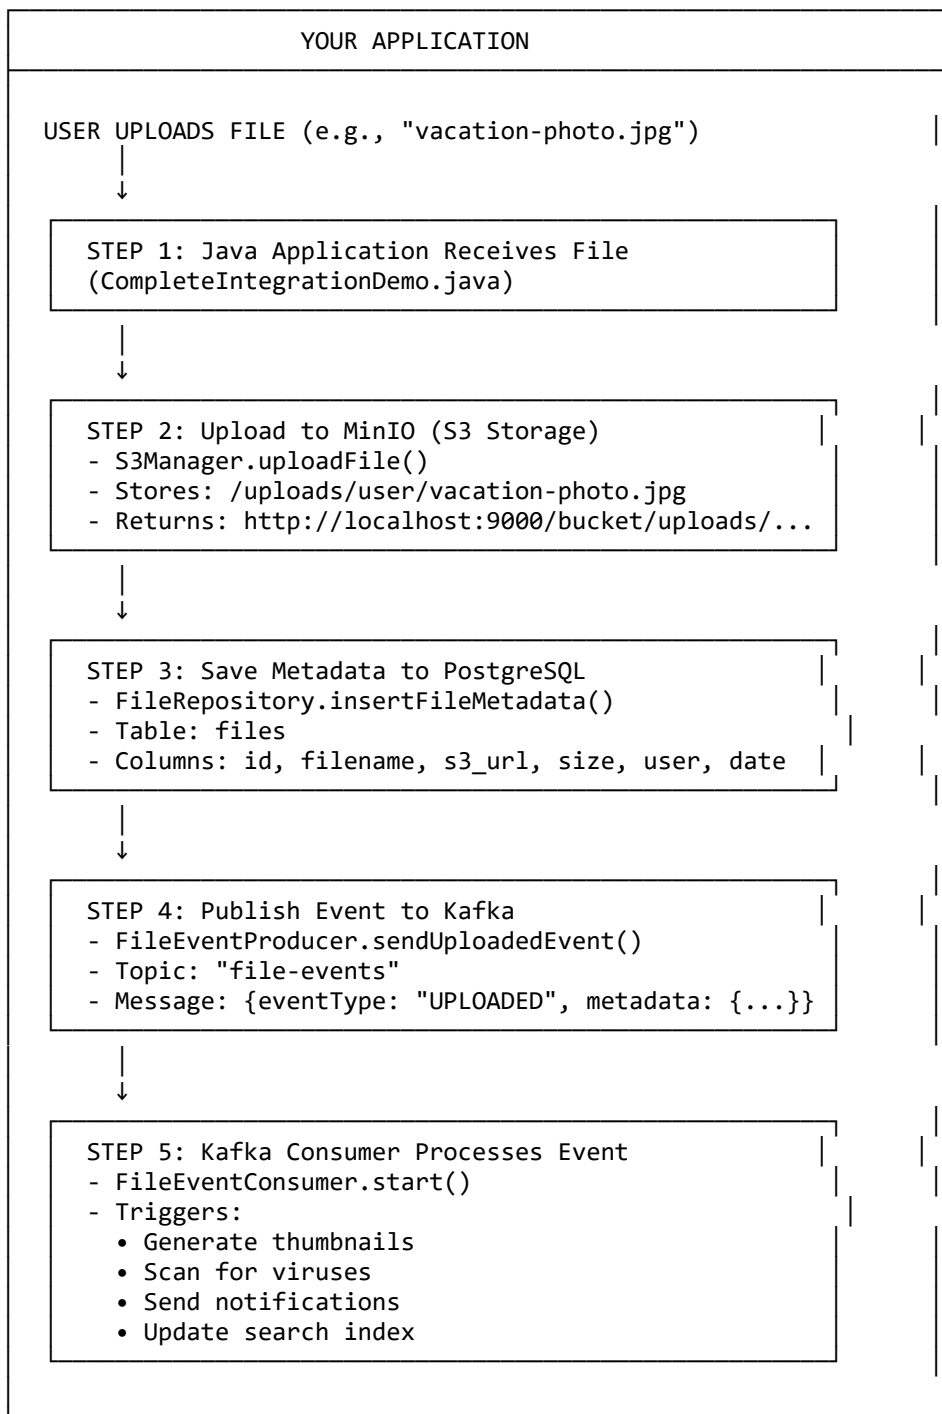
- 1. **Scalable** - Can handle millions of users
- 2. **Reliable** - If one part fails, others keep working
- 3. **Fast** - Asynchronous processing means no waiting
- 4. **Flexible** - Easy to add new features without breaking existing code
- 5. **Maintainable** - Each component has a clear responsibility

The Three Main Components

Component	Purpose	Real-World Equivalent
Apache Kafka	Message streaming & event processing	Post office that delivers messages between services
PostgreSQL	Structured data storage	Filing cabinet with organized folders
MinIO (S3)	File/blob storage	Warehouse for storing physical items

2. The Big Picture - High-Level Architecture

Visual Architecture



Why This Architecture?

Traditional Approach (Old Way):

User uploads file → Process EVERYTHING (thumbnails, scan, notify) → Response
(User waits 30 seconds!)

Event-Driven Approach (Modern Way):

User uploads file → Store file → Quick response (1 second!)
↓
Background processing happens asynchronously

The user gets immediate feedback, and heavy processing happens in the background!

3. Component 1: Apache Kafka - The Message Streaming Platform

What is Kafka?

Simple Analogy: Think of Kafka as a **super-efficient post office** that: - Never loses mail (messages) - Delivers to multiple people (consumers) simultaneously - Keeps a record of all deliveries - Can handle millions of letters per second

Why Do We Need Kafka?

Problem Without Kafka:

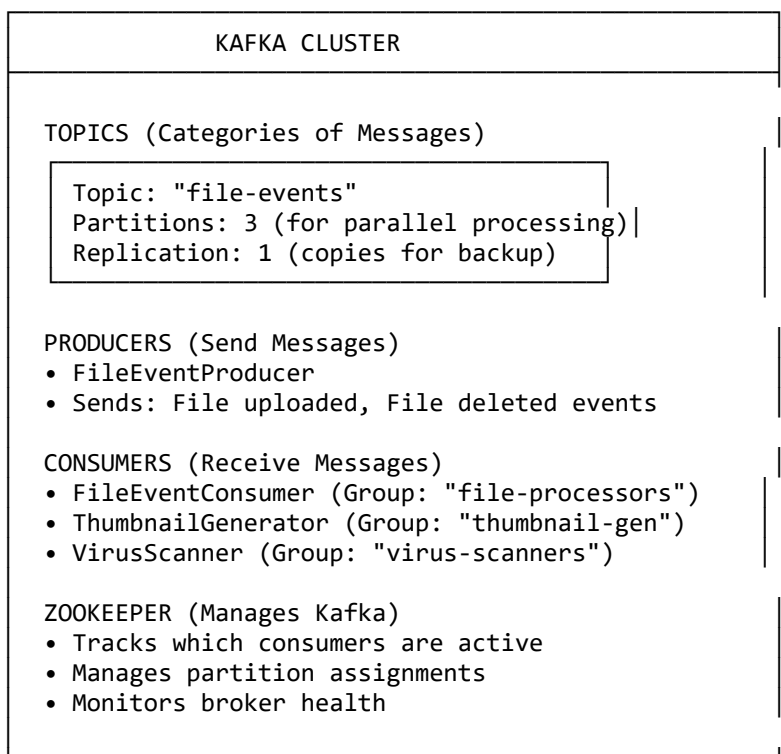
```
// Direct function calls - tightly coupled
uploadFile(file);
generateThumbnail(file); // If this fails, everything stops!
scanVirus(file);         // User waits for all of this
sendNotification(file);  // Too slow!
```

Solution With Kafka:

```
// Loosely coupled through events
uploadFile(file);
kafka.publish("file-uploaded", fileInfo); // Fire and forget!
// User gets immediate response

// Meanwhile, in the background:
// - Thumbnail service listens for "file-uploaded" events
// - Virus scanner listens for "file-uploaded" events
// - Notification service listens for "file-uploaded" events
```

Kafka Architecture



Key Kafka Concepts

1. Topics

- **What:** A category or feed name to which messages are published
- **Example:** “file-events”, “user-signups”, “payment-transactions”
- **Real-World:** Like departments in a company (HR, Finance, Engineering)

2. Partitions

- **What:** Sub-divisions of topics for parallel processing
- **Example:** Topic “file-events” has 3 partitions
 - Partition 0: Handles files from users A-H
 - Partition 1: Handles files from users I-P
 - Partition 2: Handles files from users Q-Z
- **Benefit:** 3 consumers can process simultaneously = 3x faster!

3. Producers

- **What:** Applications that send messages to Kafka topics
- **Our Code:** FileEventProducer.java
- **Example:**

```
// FileEventProducer.java:72
public void sendUploadedEvent(FileMetadata metadata) {
    FileEvent event = new FileEvent("UPLOADED", metadata);
    String jsonMessage = objectMapper.writeValueAsString(event);

    ProducerRecord<String, String> record = new ProducerRecord<>(
        topicName,           // "file-events"
        metadata.getId(),    // Partition key (keeps same user's events in order)
        jsonMessage          // The actual message
    );

    producer.send(record); // Send to Kafka
}
```

4. Consumers

- **What:** Applications that read messages from Kafka topics
- **Our Code:** FileEventConsumer.java
- **Example:**

```
// FileEventConsumer.java:87
consumer.subscribe(Collections.singletonList(topicName));

while (running) {
    ConsumerRecords<String, String> records = consumer.poll(Duration.ofMillis(100));

    for (ConsumerRecord<String, String> record : records) {
        FileEvent event = objectMapper.readValue(record.value(), FileEvent.class);

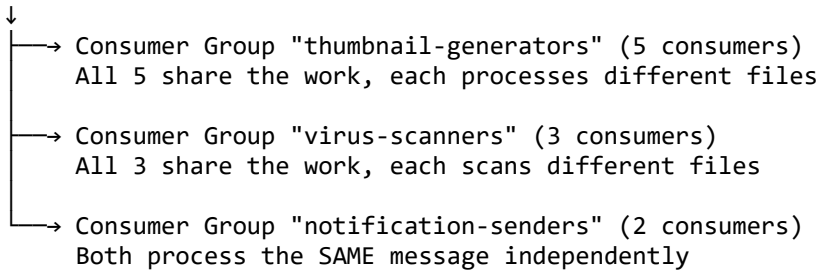
        // Process the event
        eventHandler.accept(event); // Your custom processing logic
    }
}
```

5. Consumer Groups

- **What:** Multiple consumers working together to process messages
- **Example:**

- Group “thumbnail-generators”: 5 consumers generate thumbnails in parallel
- Group “virus-scanners”: 3 consumers scan files in parallel
- Group “notification-senders”: 2 consumers send emails in parallel
- **Benefit:** Each group processes ALL messages independently

Message: "File uploaded: photo.jpg"



Kafka Configuration in Our Project

File: docker-compose.yml:21-51

```

kafka:
  image: confluentinc/cp-kafka:7.5.0
  ports:
    - "9092:9092"      # External access port
    - "29092:29092"    # Internal (Docker) access port
  environment:
    KAFKA_BROKER_ID: 1
    KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181
    KAFKA_ADVERTISED_LISTENERS:
      PLAINTEXT://kafka:29092,          # For Docker containers
      PLAINTEXT_HOST://localhost:9092  # For your Java app
    KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1
    KAFKA_AUTO_CREATE_TOPICS_ENABLE: 'true'
    KAFKA_NUM_PARTITIONS: 3            # Default partitions per topic
  
```

How Kafka Guarantees Reliability

1. **Durability:** Messages are written to disk, not just memory
2. **Replication:** Messages are copied to multiple brokers
3. **Acknowledgments:** Producer waits for confirmation
4. **Offset Tracking:** Consumers remember where they left off
5. **Retry Logic:** Failed sends are automatically retried

Kafka in Our Code

Producer Configuration (FileEventProducer.java:51):

```

Properties props = new Properties();
props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, bootstrapServers);
props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG, StringSerializer.class.getName());
props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG, StringSerializer.class.getName());
props.put(ProducerConfig.ACKS_CONFIG, "all"); // Wait for all replicas
props.put(ProducerConfig.RETRIES_CONFIG, 3); // Retry 3 times
props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, "true"); // Prevent duplicates
  
```

Consumer Configuration (FileEventConsumer.java:58):

```

Properties props = new Properties();
props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, bootstrapServers);
props.put(ConsumerConfig.GROUP_ID_CONFIG, groupId);
props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG, StringDeserializer.class.getName());
props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG, StringDeserializer.class.getName());
  
```

```
props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest"); // Read from beginning
props.put(ConsumerConfig.ENABLE_AUTO_COMMIT_CONFIG, "true"); // Auto-save progress
```

4. Component 2: PostgreSQL - The Relational Database

What is PostgreSQL?

Simple Analogy: Think of PostgreSQL as a **highly organized filing cabinet** where: - Each drawer is a **table** - Each file is a **row** - Each label on the file is a **column** - The filing system ensures everything is organized and easy to find

Why Do We Need a Database?

What We Store in PostgreSQL: - **Metadata:** Information ABOUT files (not the files themselves) - **Structured Data:** Data with clear relationships - **Searchable Information:** Things you need to query quickly

What We DON'T Store in PostgreSQL: - Large files (photos, videos, documents) → Goes to S3/MinIO - Temporary data → Goes to Redis/Memcached - Real-time events → Goes to Kafka

Database Schema

File: init-db.sql:6-18

```
-- Table to store file metadata
CREATE TABLE IF NOT EXISTS files (
    id VARCHAR(255) PRIMARY KEY,           -- Unique identifier (UUID)
    filename VARCHAR(500) NOT NULL,        -- Original filename
    s3_key VARCHAR(1000) NOT NULL,         -- Path in S3 bucket
    s3_url TEXT NOT NULL,                  -- Full URL to access file
    file_size BIGINT NOT NULL,             -- Size in bytes
    content_type VARCHAR(255),             -- MIME type (image/jpeg, video/mp4)
    uploaded_by VARCHAR(255) NOT NULL,     -- Username of uploader
    uploaded_at BIGINT NOT NULL,           -- Unix timestamp
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP, -- Auto-populated
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP -- Auto-updated
);
```

What Each Column Means:

Column	Type	Purpose	Example
id	VARCHAR(255)	Unique identifier	“a1b2c3d4-e5f6-7g8h-9i0j-k1l2m3n4o5p6”
filename	VARCHAR(500)	Original name	“vacation-photo.jpg”
s3_key	VARCHAR(1000)	S3 storage path	“uploads/alice/vacation-photo.jpg”
s3_url	TEXT	Complete URL	“http://localhost:9000/bucket/uploads/alice/vacation-photo.jpg”
file_size	BIGINT	Size in bytes	2621440 (2.5 MB)
content_type	VARCHAR(255)	File type	“image/jpeg”
uploaded_by	VARCHAR(255)	User	“alice”
uploaded_at	BIGINT	Unix timestamp	1729090624308
created_at	TIMESTAMP	Row creation time	“2025-10-16 16:37:04”
updated_at	TIMESTAMP	Last update time	“2025-10-16 16:37:04”

Database Indexes

File: init-db.sql:21-24

```
-- Indexes for faster queries
CREATE INDEX IF NOT EXISTS idx_files_uploaded_by ON files(uploaded_by);
CREATE INDEX IF NOT EXISTS idx_files_created_at ON files(created_at DESC);
CREATE INDEX IF NOT EXISTS idx_files_content_type ON files(content_type);
```

Why Indexes?

Without Index:

```
SELECT * FROM files WHERE uploaded_by = 'alice';
-- Database scans ALL 1,000,000 rows → Takes 5 seconds!
```

With Index:

```
SELECT * FROM files WHERE uploaded_by = 'alice';
-- Database uses index → Finds 10 rows instantly → Takes 0.01 seconds!
```

Analogy: Like the index at the back of a textbook - instead of reading every page, you jump directly to the relevant page!

Connection Pooling with HikariCP

What is Connection Pooling?

Without Pooling (Slow):

Request 1: Create DB connection → Use it → Close it
 Request 2: Create DB connection → Use it → Close it
 Request 3: Create DB connection → Use it → Close it
 (Each connection creation takes 100ms!)

With Pooling (Fast):

Startup: Create 10 connections in advance (1 second total)
 Request 1: Borrow connection 1 → Use it → Return it
 Request 2: Borrow connection 2 → Use it → Return it
 Request 3: Borrow connection 1 → Use it → Return it
 (Each request takes 0.1ms!)

Our Configuration (DatabaseManager.java:30-50):

```
HikariConfig config = new HikariConfig();
config.setJdbcUrl("jdbc:postgresql://localhost:5432/mydb");
config.setUsername("admin");
config.setPassword("admin123");

// Pool settings
config.setMaximumPoolSize(10);           // Max 10 connections
config.setMinimumIdle(2);                 // Always keep 2 ready
config.setConnectionTimeout(30000);       // Wait max 30s for connection
config.setIdleTimeout(600000);            // Close idle connections after 10 min
config.setMaxLifetime(1800000);           // Renew connections every 30 min

// Performance optimizations
config.addDataSourceProperty("cachePrepStmts", "true");
config.addDataSourceProperty("prepStmtCacheSize", "250");
config.addDataSourceProperty("prepStmtCacheSqlLimit", "2048");

dataSource = new HikariDataSource(config);
```

CRUD Operations

File Repository (FileRepository.java)

Create (Insert)

```
// FileRepository.java:48-76
public void insertFileMetadata(FileMetadata metadata) throws SQLException {
    String sql = ""
        INSERT INTO files (id, filename, s3_key, s3_url, file_size,
                           content_type, uploaded_by, uploaded_at)
        VALUES (?, ?, ?, ?, ?, ?, ?, ?)
        "";

    try (Connection conn = databaseManager.getConnection();
        PreparedStatement stmt = conn.prepareStatement(sql)) {

        stmt.setString(1, metadata.getId());
        stmt.setString(2, metadata.getFilename());
        stmt.setString(3, metadata.getS3Key());
        stmt.setString(4, metadata.getS3Url());
        stmt.setLong(5, metadata.getFileSize());
        stmt.setString(6, metadata.getContentType());
        stmt.setString(7, metadata.getUploadedBy());
        stmt.setLong(8, metadata.getUploadedAt());

        stmt.executeUpdate();
        logger.info("Inserted file metadata: {}", metadata.getFilename());
    }
}
```

Read (Query)

```
// FileRepository.java:83-108
public FileMetadata getFileById(String id) throws SQLException {
    String sql = "SELECT * FROM files WHERE id = ?";

    try (Connection conn = databaseManager.getConnection();
        PreparedStatement stmt = conn.prepareStatement(sql)) {

        stmt.setString(1, id);
        ResultSet rs = stmt.executeQuery();

        if (rs.next()) {
            return new FileMetadata(
                rs.getString("id"),
                rs.getString("filename"),
                rs.getString("s3_key"),
                rs.getString("s3_url"),
                rs.getLong("file_size"),
                rs.getString("content_type"),
                rs.getString("uploaded_by"),
                rs.getLong("uploaded_at")
            );
        }
        return null;
    }
}
```

Update

```
// FileRepository.java:115-132
public void updateFileMetadata(String id, FileMetadata metadata) throws SQLException {
    String sql = ""
```



```

UPDATE files
SET filename = ?, s3_key = ?, s3_url = ?, file_size = ?,
    content_type = ?, updated_at = CURRENT_TIMESTAMP
WHERE id = ?
""";

try (Connection conn = databaseManager.getConnection();
    PreparedStatement stmt = conn.prepareStatement(sql)) {

    stmt.setString(1, metadata.getFilename());
    stmt.setString(2, metadata.getS3Key());
    stmt.setString(3, metadata.getS3Url());
    stmt.setLong(4, metadata.getFileSize());
    stmt.setString(5, metadata.getContentType());
    stmt.setString(6, id);

    stmt.executeUpdate();
}
}

```

Delete

```

// FileRepository.java:139-153
public void deleteFileMetadata(String id) throws SQLException {
    String sql = "DELETE FROM files WHERE id = ?";

    try (Connection conn = databaseManager.getConnection();
        PreparedStatement stmt = conn.prepareStatement(sql)) {

        stmt.setString(1, id);
        int rowsAffected = stmt.executeUpdate();

        if (rowsAffected > 0) {
            logger.info("Deleted file metadata: {}", id);
        }
    }
}

```

Transactions

What is a Transaction?

A transaction ensures that **all operations succeed together, or all fail together**.

Example Without Transaction:

```

// BAD: If step 2 fails, step 1 is already done!
updateUserBalance(userId, -100);    // Deduct $100
updateMerchantBalance(merchantId, +100); // This fails!
// User lost $100, merchant didn't get it! ❌

```

Example With Transaction:

```

// GOOD: Both succeed or both fail
try (Connection conn = dataSource.getConnection()) {
    conn.setAutoCommit(false); // Start transaction

    updateUserBalance(conn, userId, -100);
    updateMerchantBalance(conn, merchantId, +100);

    conn.commit(); // Both succeeded! ✅
} catch (Exception e) {

```

```
conn.rollback(); // Undo everything! 
}
```

PostgreSQL Configuration

File: docker-compose.yml:89-104

```
postgres:
  image: postgres:16-alpine
  container_name: postgres-db
  environment:
    POSTGRES_USER: admin          # Username
    POSTGRES_PASSWORD: admin123  # Password
    POSTGRES_DB: mydb            # Database name
  ports:
    - "5432:5432"                # Port mapping
  volumes:
    - postgres-data:/var/lib/postgresql/data # Persistent storage
  healthcheck:
    test: ["CMD-SHELL", "pg_isready -U admin -d mydb"]
    interval: 10s
    timeout: 5s
    retries: 5
```

Common SQL Queries

Query 1: Get all files uploaded by a user

```
SELECT * FROM files
WHERE uploaded_by = 'alice'
ORDER BY created_at DESC
LIMIT 10;
```

Query 2: Get total storage used by user

```
SELECT uploaded_by, SUM(file_size) as total_bytes
FROM files
GROUP BY uploaded_by;
```

Query 3: Find large files

```
SELECT filename, file_size, uploaded_by
FROM files
WHERE file_size > 10485760 -- 10 MB
ORDER BY file_size DESC;
```

Query 4: Get file statistics by type

```
SELECT
  content_type,
  COUNT(*) as file_count,
  SUM(file_size) as total_size,
  AVG(file_size) as avg_size
FROM files
GROUP BY content_type;
```

5. Component 3: MinIO (S3) - The Blob Storage

What is MinIO?

Simple Analogy: Think of MinIO as a **massive warehouse** where: - Each item (file) has a unique barcode (S3 key) - Items are organized in sections (buckets) - You can retrieve any item instantly using its barcode - The warehouse is virtually unlimited in size

MinIO vs AWS S3: - **AWS S3:** Cloud storage service by Amazon - **MinIO:** Open-source, self-hosted, S3-compatible storage - **API:** MinIO uses the EXACT same API as AWS S3 - **Benefit:** Develop locally with MinIO, deploy to AWS S3 with zero code changes!

Why Separate File Storage from Database?

Storing Files in Database (Bad Idea ❌):

Database size: 100 GB

- User data: 1 GB
- Files: 99 GB (images, videos, documents)

Problems:

- ❌ Slow queries (database scans huge files)
- ❌ Expensive (database storage costs 10x more)
- ❌ Backups take forever
- ❌ Can't scale independently

Storing Files in S3/MinIO (Good Idea ✅):

Database: 1 GB (just metadata)

S3 Storage: 99 GB (actual files)

Benefits:

- ✅ Fast queries (database only scans metadata)
- ✅ Cheap (S3 costs 10x less than database storage)
- ✅ Quick backups
- ✅ Scales independently

S3 Architecture Concepts

1. Buckets

- **What:** Top-level containers for storing objects
- **Example:** “user-uploads”, “profile-pictures”, “video-content”
- **Analogy:** Like a hard drive or USB stick
- **Our Bucket:** “complete-integration-bucket”

2. Objects (Files)

- **What:** The actual files you store
- **Components:**
 - **Key:** Unique identifier/path (e.g., “uploads/alice/photo.jpg”)
 - **Data:** The file content
 - **Metadata:** Additional info (content-type, size, etc.)

3. Keys (File Paths)

- **What:** The “address” of an object in a bucket
- **Format:** Like a file system path
- **Examples:**
 - uploads/alice/vacation-photo.jpg
 - documents/2025/report.pdf
 - videos/tutorial-01.mp4

4. URLs

- **What:** Web address to access the file
- **Format:** http://endpoint/bucket-name/key
- **Example:** http://localhost:9000/complete-integration-bucket/uploads/alice/photo.jpg

MinIO Configuration

File: docker-compose.yml:120-136

```
minio:
  image: minio/minio:latest
  container_name: minio-s3
  command: server /data --console-address ":9001"
  environment:
    MINIO_ROOT_USER: minioadmin      # Access key
    MINIO_ROOT_PASSWORD: minioadmin123 # Secret key
  ports:
    - "9000:9000"      # API port (S3 operations)
    - "9001:9001"      # Web console port
  volumes:
    - minio-data:/data # Persistent storage
  healthcheck:
    test: ["CMD", "curl", "-f", "http://localhost:9000/minio/health/live"]
    interval: 10s
    timeout: 5s
    retries: 5
```

S3Manager - Java Implementation

File: S3Manager.java

Initialization

```
// S3Manager.java:33-55
public S3Manager(String endpoint, String accessKey, String secretKey) {
    // Configure AWS SDK for MinIO
    AwsBasicCredentials credentials = AwsBasicCredentials.create(accessKey, secretKey);

    this.s3Client = S3Client.builder()
        .endpointOverride(URI.create(endpoint)) // Point to MinIO instead of AWS
        .region(Region.US_EAST_1)
        .credentialsProvider(StaticCredentialsProvider.create(credentials))
        .serviceConfiguration(S3Configuration.builder()
            .pathStyleAccessEnabled(true) // Required for MinIO
            .build())
        .build();
}
```

Create Bucket

```
// S3Manager.java:62-77
public void createBucketIfNotExists(String bucketName) {
    try {
        // Check if bucket exists
        s3Client.headBucket(HeadBucketRequest.builder()
            .bucket(bucketName)
            .build());

        logger.info("Bucket {} already exists", bucketName);
    }
```

```

    } catch (NoSuchBucketException e) {
        // Bucket doesn't exist, create it
        s3Client.createBucket(CreateBucketRequest.builder()
            .bucket(bucketName)
            .build());

        logger.info("Created bucket: {}", bucketName);
    }
}

```

Upload File

```

// S3Manager.java:90-119
public String uploadFile(String bucketName, String key, byte[] fileData, String contentType) {
    try {
        // Upload file to S3
        PutObjectRequest request = PutObjectRequest.builder()
            .bucket(bucketName)
            .key(key)
            .contentType(contentType)
            .contentLength((long) fileData.length)
            .build();

        s3Client.putObject(request, RequestBody.fromBytes(fileData));

        // Generate URL
        String url = String.format("%s/%s/%s",
            s3Client.serviceClientConfiguration().endpointOverride().get(),
            bucketName,
            key);

        logger.info("Uploaded file to S3: {} ({})", key, formatFileSize(fileData.length));
        return url;
    } catch (Exception e) {
        logger.error("Failed to upload file to S3: {}", key, e);
        throw new RuntimeException("S3 upload failed", e);
    }
}

```

What Happens When You Upload: 1. File is divided into chunks (for large files) 2. Each chunk is uploaded separately 3. MinIO reassembles chunks on the server 4. Returns a URL to access the file 5. File is immediately available worldwide

Download File

```

// S3Manager.java:128-147
public byte[] downloadFile(String bucketName, String key) {
    try {
        GetObjectRequest request = GetObjectRequest.builder()
            .bucket(bucketName)
            .key(key)
            .build();

        ResponseInputStream<GetObjectResponse> response = s3Client.getObject(request);
        byte[] fileData = response.readAllBytes();

        logger.info("Downloaded file from S3: {} ({})", key, formatFileSize(fileData.length));
        return fileData;
    } catch (NoSuchKeyException e) {
        logger.error("File not found in S3: {}", key);
        throw new RuntimeException("File not found", e);
    }
}

```

```
    }
}
```

Delete File

```
// S3Manager.java:154-169
public void deleteFile(String bucketName, String key) {
    try {
        DeleteObjectRequest request = DeleteObjectRequest.builder()
            .bucket(bucketName)
            .key(key)
            .build();

        s3Client.deleteObject(request);
        logger.info("Deleted file from S3: {}", key);

    } catch (Exception e) {
        logger.error("Failed to delete file from S3: {}", key, e);
        throw new RuntimeException("S3 delete failed", e);
    }
}
```

List Files

```
// S3Manager.java:176-201
public List<String> listFiles(String bucketName, String prefix) {
    List<String> fileKeys = new ArrayList<>();

    try {
        ListObjectsV2Request request = ListObjectsV2Request.builder()
            .bucket(bucketName)
            .prefix(prefix)
            .build();

        ListObjectsV2Response response = s3Client.listObjectsV2(request);

        for (S3Object object : response.contents()) {
            fileKeys.add(object.key());
        }

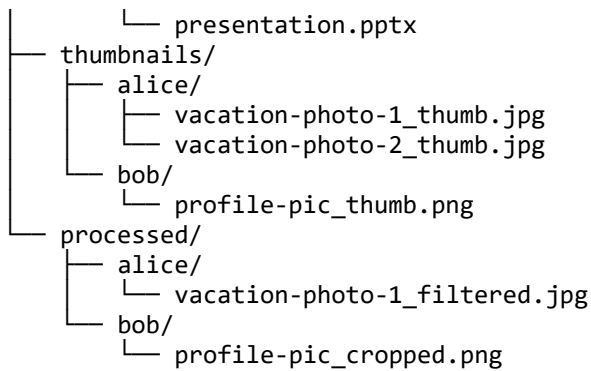
        logger.info("Listed {} files with prefix: {}", fileKeys.size(), prefix);
        return fileKeys;

    } catch (Exception e) {
        logger.error("Failed to list files from S3", e);
        throw new RuntimeException("S3 list failed", e);
    }
}
```

File Organization Strategy

Good File Organization:

```
bucket: complete-integration-bucket
├── uploads/
│   ├── alice/
│   │   ├── vacation-photo-1.jpg
│   │   ├── vacation-photo-2.jpg
│   │   └── document.pdf
│   ├── bob/
│   │   ├── profile-pic.png
│   │   └── resume.pdf
│   └── carol/
```



Benefits: - Easy to find user's files - Can set permissions per folder - Simple to backup specific users - Clear separation of file types

S3 Best Practices

1. Unique Keys

```

// BAD: Filename collisions
String key = "uploads/" + filename; // Multiple users upload "photo.jpg"

// GOOD: Use UUID or timestamp
String key = "uploads/" + userId + "/" + UUID.randomUUID() + "_" + filename;

```

2. Content-Type

```

// Set proper content-type for browsers to handle correctly
String contentType = "image/jpeg"; // Browser displays image
String contentType = "video/mp4"; // Browser plays video
String contentType = "application/pdf"; // Browser shows PDF

```

3. File Size Limits

```

// Check file size before upload
long MAX_FILE_SIZE = 100 * 1024 * 1024; // 100 MB

if (fileData.length > MAX_FILE_SIZE) {
    throw new IllegalArgumentException("File too large!");
}

```

4. Error Handling

```

try {
    s3Manager.uploadFile(bucket, key, data, contentType);
} catch (S3Exception e) {
    if (e.getStatusCode() == 403) {
        // Permission denied
    } else if (e.getStatusCode() == 404) {
        // Bucket not found
    } else if (e.getStatusCode() == 503) {
        // Service unavailable, retry
    }
}

```

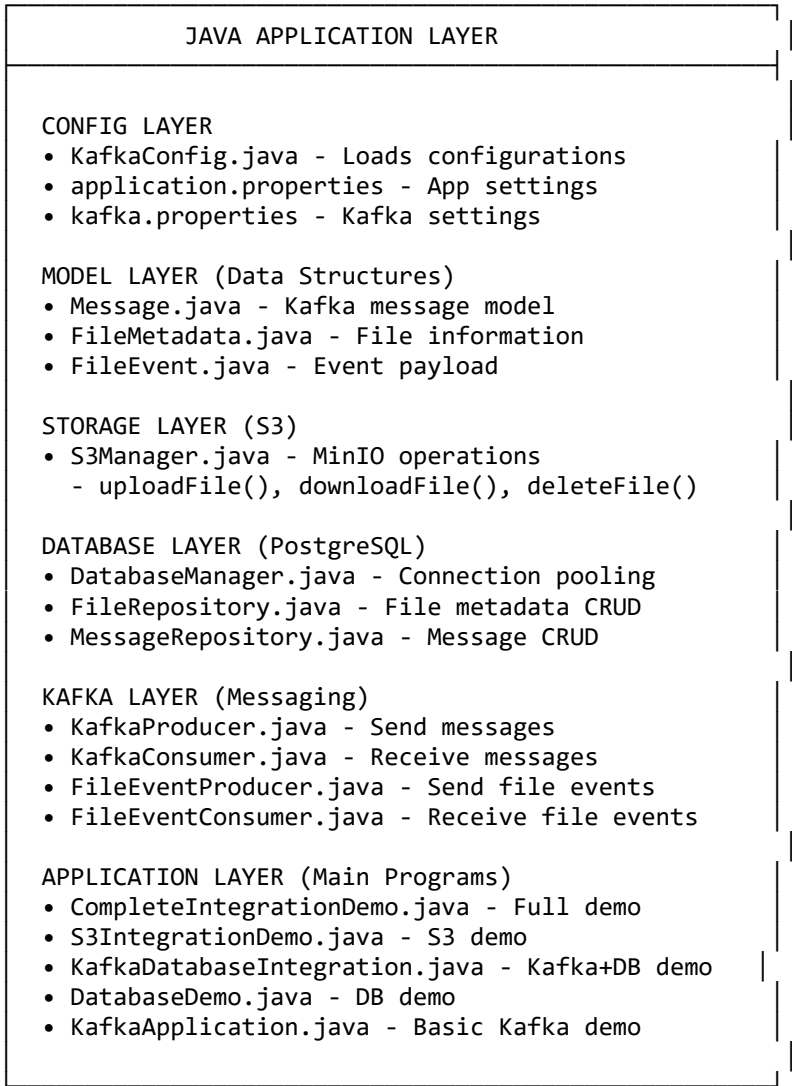
S3 vs Database: When to Use What?

Data Type	Store In	Reason
Images	S3	Large, binary, served directly to users

Data Type	Store In	Reason
Videos	S3	Very large, streaming required
Documents (PDF, DOCX)	S3	Large, downloadable
User profiles	Database	Small, structured, frequently queried
Comments	Database	Small, structured, searchable
Thumbnails	S3	Images, but smaller versions
Log files	S3	Large, infrequently accessed
Session data	Redis/Cache	Temporary, needs fast access
Configurations	Database	Small, structured, version controlled

6. Component 4: Java Application - The Orchestrator

Application Architecture



Complete Integration Flow

File: CompleteIntegrationDemo.java

Main Method

```
// CompleteIntegrationDemo.java:56-78
public static void main(String[] args) {
    logger.info("=".repeat(70));
```



```

logger.info(" COMPLETE S3 + KAFKA + DATABASE INTEGRATION DEMO");
logger.info("=".repeat(70));

CompleteIntegrationDemo demo = new CompleteIntegrationDemo();

try {
    // Initialize all components
    demo.initializeComponents();

    // Start consumer in background
    demo.startConsumer();

    // Upload sample files
    demo.uploadSampleFiles();

    // Wait for processing
    Thread.sleep(5000);

    // Show statistics
    demo.displayStatistics();

} finally {
    demo.cleanup();
}
}

```

Component Initialization

```

// CompleteIntegrationDemo.java:85-115
private void initializeComponents() {
    logger.info("\n" + "=".repeat(70));
    logger.info(" INITIALIZING COMPONENTS");
    logger.info("=".repeat(70));

    // 1. Initialize S3 (MinIO)
    logger.info("\n[1/4] Initializing S3 Storage (MinIO)...");
    s3Manager = new S3Manager(
        "http://localhost:9000",
        "minioadmin",
        "minioadmin123"
    );
    s3Manager.createBucketIfNotExists(BUCKET_NAME);

    // 2. Initialize Database
    logger.info("[2/4] Initializing PostgreSQL Database...");
    databaseManager = new DatabaseManager(
        "jdbc:postgresql://localhost:5432/mydb",
        "admin",
        "admin123"
    );
    fileRepository = new FileRepository(databaseManager);

    // 3. Initialize Kafka Producer
    logger.info("[3/4] Initializing Kafka Producer...");
    eventProducer = new FileEventProducer(
        "localhost:9092",
        "file-events"
    );

    // 4. Initialize Kafka Consumer
    logger.info("[4/4] Initializing Kafka Consumer...");
    eventConsumer = new FileEventConsumer(
        "localhost:9092",
        "complete-integration-group",

```

```

        "file-events"
    );

    logger.info("\n✅ All components initialized successfully!\n");
}

```

File Upload Process

```

// CompleteIntegrationDemo.java:156-214
private void uploadFile(String filename, String contentType) {
    try {
        logger.info("\n" + "-".repeat(70));
        logger.info("📁 UPLOADING: {}", filename);
        logger.info("-".repeat(70));

        // STEP 1: Create sample file data
        byte[] fileData = ("Sample content for " + filename).getBytes();
        String fileId = UUID.randomUUID().toString();
        String s3Key = "uploads/" + DEMO_USER + "/" + filename;

        logger.info(" [1/4] Created file: {} ({} bytes)", filename, fileData.length);

        // STEP 2: Upload to S3
        logger.info(" [2/4] Uploading to S3...");
        String s3Url = s3Manager.uploadFile(BUCKET_NAME, s3Key, fileData, contentType);
        logger.info("      ✅ S3 URL: {}", s3Url);

        // STEP 3: Create metadata object
        FileMetadata metadata = new FileMetadata(
            fileId,
            filename,
            s3Key,
            s3Url,
            (long) fileData.length,
            contentType,
            DEMO_USER,
            System.currentTimeMillis()
        );

        // STEP 4: Save metadata to database
        logger.info(" [3/4] Saving metadata to PostgreSQL...");
        fileRepository.insertFileMetadata(metadata);
        logger.info("      ✅ Saved to database with ID: {}", fileId);

        // STEP 5: Send event to Kafka
        logger.info(" [4/4] Publishing event to Kafka...");
        eventProducer.sendUploadedEvent(metadata);
        logger.info("      ✅ Event published to topic: file-events");

        logger.info("\n✅ File upload complete: {}", filename);
        filesUploaded.incrementAndGet();

    } catch (Exception e) {
        logger.error("❌ Failed to upload file: {}", filename, e);
    }
}

```

Consumer Event Handler

```

// CompleteIntegrationDemo.java:125-147
private void startConsumer() {
    logger.info("\n" + "=".repeat(70));
    logger.info(" STARTING EVENT CONSUMER");
}

```

```

logger.info("=".repeat(70));

// Define event handler
Consumer<FileEvent> eventHandler = event -> {
    try {
        logger.info("\n" + "=".repeat(70));
        logger.info("📁 RECEIVED EVENT: {}", event.getEventType());
        logger.info("=".repeat(70));

        FileMetadata metadata = event.getFileMetadata();
        logger.info(" • Filename: {}", metadata.getFilename());
        logger.info(" • Size: {}", formatFileSize(metadata.getFileSize()));
        logger.info(" • User: {}", metadata.getUploadedBy());
        logger.info(" • S3 URL: {}", metadata.getS3Url());

        // Simulate processing
        logger.info("\n 🔄 Processing tasks:");
        logger.info("    [1] Generating thumbnails...");
        logger.info("    [2] Scanning for viruses...");
        logger.info("    [3] Updating search index...");
        logger.info("    [4] Sending notifications...");

        Thread.sleep(100); // Simulate work

        logger.info("\n ✅ Processing complete!");
        eventsProcessed.incrementAndGet();

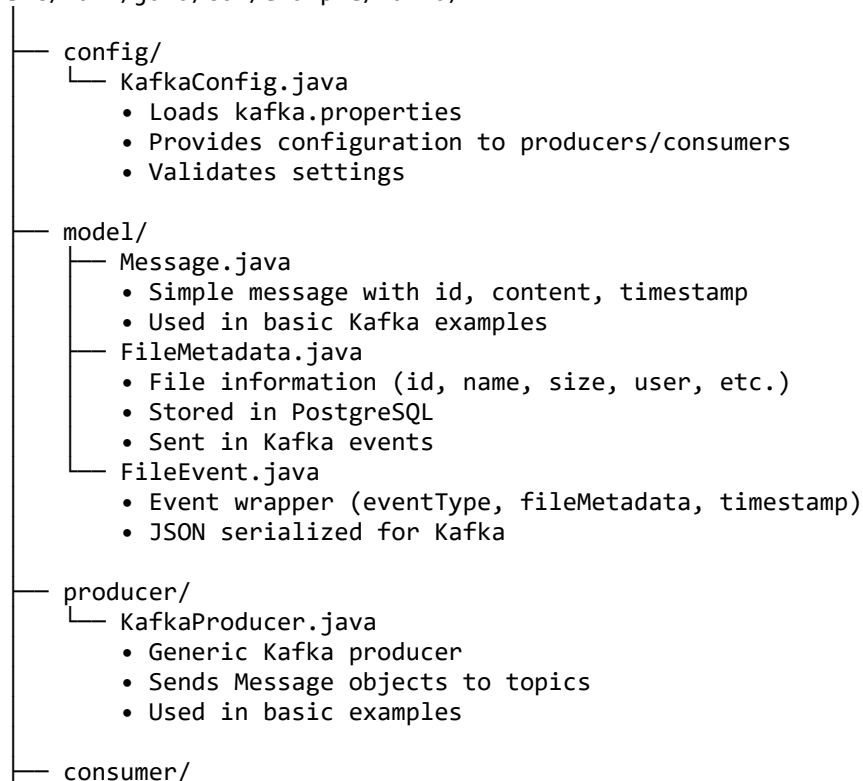
    } catch (Exception e) {
        logger.error("❌ Error processing event", e);
    }
};

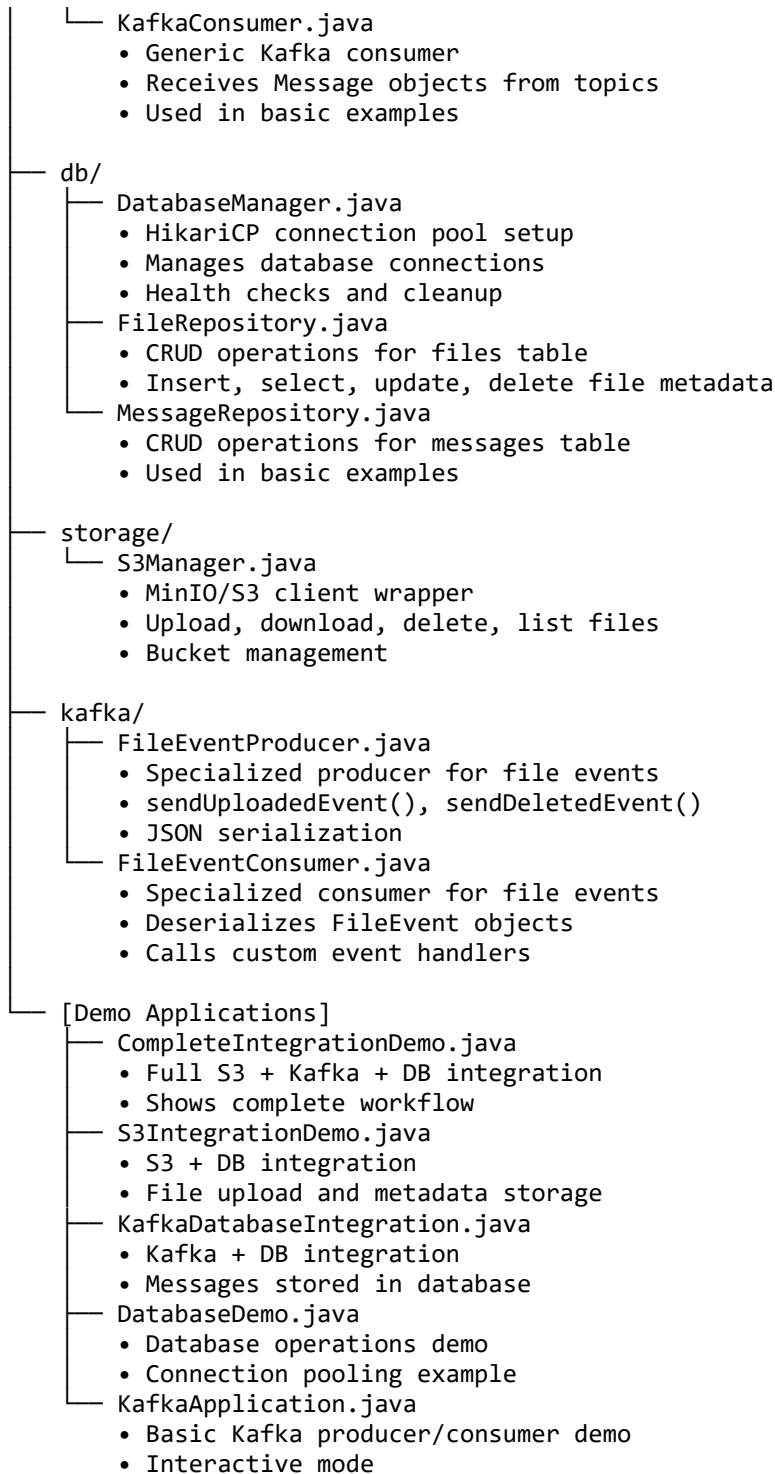
// Start consumer in background thread
new Thread(() -> eventConsumer.start(eventHandler)).start();
logger.info("✅ Consumer started and listening for events...\n");
}

```

Project Structure Explained

src/main/java/com/example/kafka/





Maven Dependencies

File: pom.xml

```

<dependencies>
  <!-- Kafka Client -->
  <dependency>
    <groupId>org.apache.kafka</groupId>
    <artifactId>kafka-clients</artifactId>
    <version>3.6.0</version>
  </dependency>

  <!-- Logging -->
  <dependency>
    <groupId>org.slf4j</groupId>
    <artifactId>slf4j-api</artifactId>
    <version>2.0.9</version>
  </dependency>

```

```

</dependency>
<dependency>
  <groupId>ch.qos.logback</groupId>
  <artifactId>logback-classic</artifactId>
  <version>1.4.11</version>
</dependency>

<!-- JSON Processing -->
<dependency>
  <groupId>com.fasterxml.jackson.core</groupId>
  <artifactId>jackson-databind</artifactId>
  <version>2.16.0</version>
</dependency>

<!-- PostgreSQL Driver -->
<dependency>
  <groupId>org.postgresql</groupId>
  <artifactId>postgresql</artifactId>
  <version>42.7.1</version>
</dependency>

<!-- Connection Pooling -->
<dependency>
  <groupId>com.zaxxer</groupId>
  <artifactId>HikariCP</artifactId>
  <version>5.1.0</version>
</dependency>

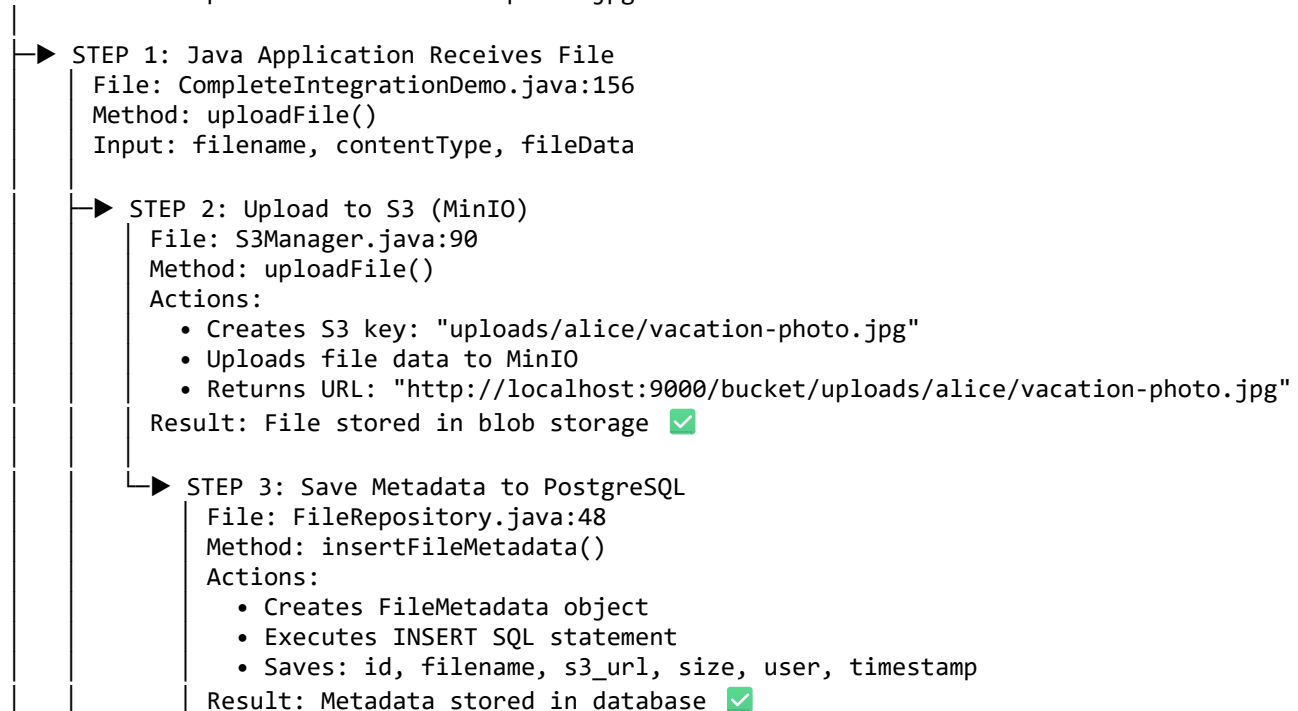
<!-- AWS S3 SDK (works with MinIO) -->
<dependency>
  <groupId>software.amazon.awssdk</groupId>
  <artifactId>s3</artifactId>
  <version>2.21.0</version>
</dependency>
</dependencies>

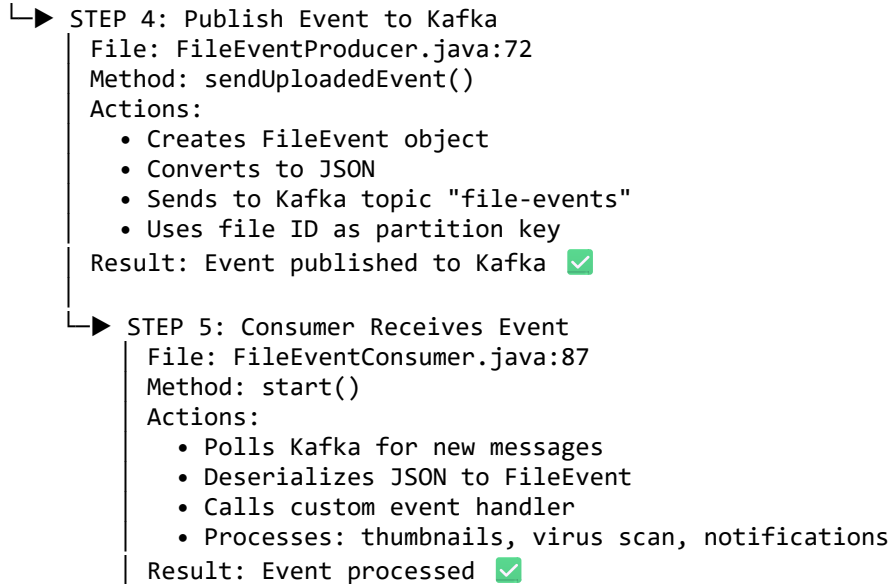
```

7. How Everything Works Together

The Complete Workflow

USER ACTION: Upload file "vacation-photo.jpg"

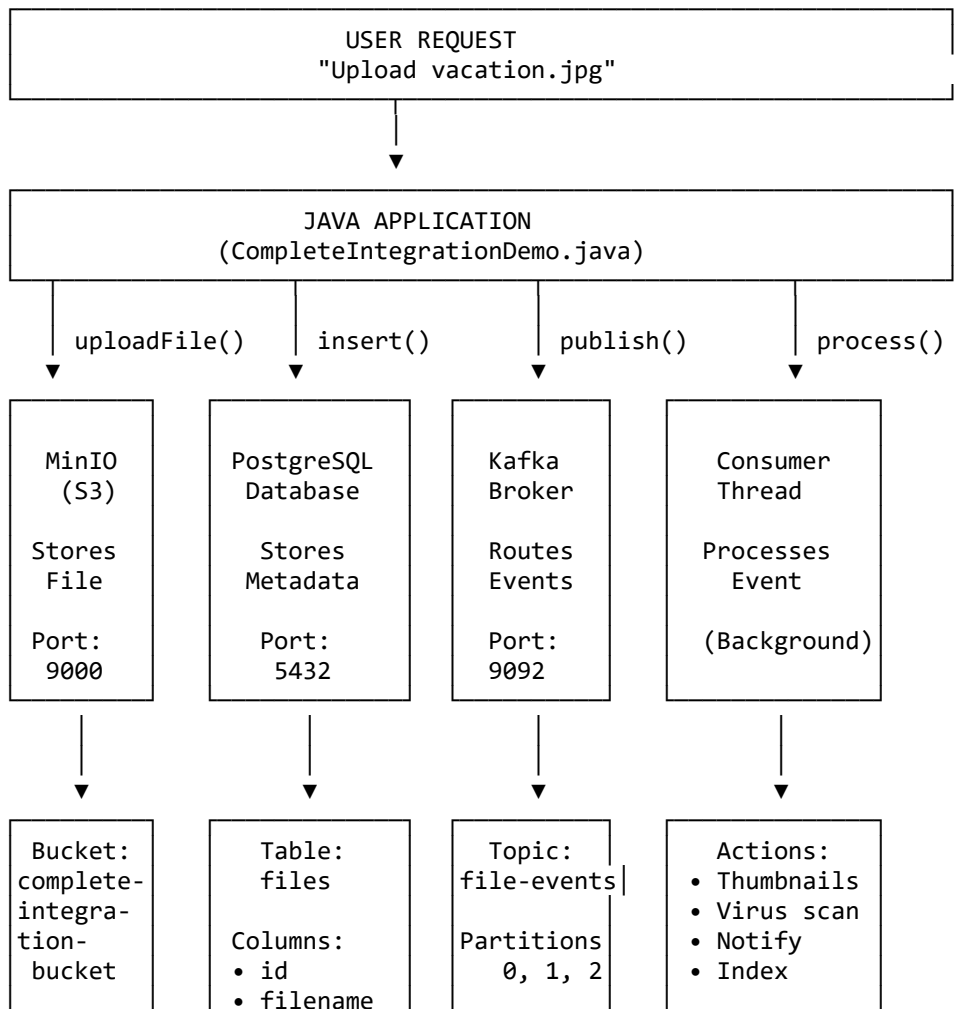




RESPONSE: Return success to user (1 second)
User sees: "File uploaded successfully!"
Background: Processing continues asynchronously

BACKGROUND PROCESSING (runs independently)
• Thumbnail generation (10 seconds)
• Virus scanning (15 seconds)
• Search indexing (5 seconds)
• Email notification (2 seconds)
Total: 32 seconds (but user doesn't wait!)

Component Interaction Diagram



Files: /uploads /alice /*.jpg	<ul style="list-style-type: none">• s3_url• size• user• timestamp	Messages: JSON Events	Async Processing
--	--	-----------------------------	---------------------

Data Flow Timeline

Time	Component	Action
0.0s	User	Clicks "Upload" button
0.1s	Java App	Receives file (10 MB)
0.2s	S3Manager	Starts upload to MinIO
1.0s	MinIO	File stored, returns URL
1.1s	FileRepository	Saves metadata to PostgreSQL
1.2s	PostgreSQL	INSERT complete
1.3s	FileEventProducer	Creates JSON event
1.4s	Kafka Broker	Receives event, stores in partition 1
1.5s	Java App	Returns success to user
		[User sees "Upload successful!"]
2.0s	Consumer Thread	Polls Kafka, receives event
2.1s	Event Handler	Starts processing
2.2s	Thumbnail Service	Generates 3 sizes (200x200, 500x500, 1000x1000)
12.0s	Thumbnail Service	Complete
12.1s	Virus Scanner	Scans file for malware
27.0s	Virus Scanner	Complete (Clean)
27.1s	Search Indexer	Extracts metadata, updates search DB
32.0s	Search Indexer	Complete
32.1s	Notification	Sends email: "Your file is ready!"
34.0s	Notification	Complete
	TOTAL USER WAIT	1.5 seconds
	TOTAL PROCESSING	34.0 seconds (async)

Error Handling & Resilience

Scenario 1: S3 Upload Fails

```
try {
    s3Url = s3Manager.uploadFile(bucket, key, data, contentType);
} catch (S3Exception e) {
    logger.error("S3 upload failed: {}", e.getMessage());
    // Don't save to database or send Kafka event
    return "Upload failed - please try again";
}
```

Result: User sees error, database and Kafka stay clean

Scenario 2: Database Insert Fails

```
try {
    fileRepository.insertFileMetadata(metadata);
} catch (SQLException e) {
    logger.error("Database insert failed: {}", e.getMessage());
    // Rollback: Delete from S3
    s3Manager.deleteFile(bucket, key);
    return "Upload failed - please try again";
}
```

Result: S3 file deleted, user sees error, Kafka not notified

Scenario 3: Kafka Publish Fails

```

try {
    eventProducer.sendUploadedEvent(metadata);
} catch (Exception e) {
    logger.error("Kafka publish failed: {}", e.getMessage());
    // File is uploaded and in database
    // Event will be retried later or processing triggered manually
}

```

Result: File uploaded and saved, but processing delayed

Scenario 4: Consumer Processing Fails

```

// FileEventConsumer.java
consumer.subscribe(Collections.singletonList(topicName));

while (running) {
    try {
        ConsumerRecords<String, String> records = consumer.poll(Duration.ofMillis(100));

        for (ConsumerRecord<String, String> record : records) {
            try {
                FileEvent event = deserialize(record.value());
                eventHandler.accept(event); // Process event

            } catch (Exception e) {
                logger.error("Failed to process event: {}", record.key(), e);
                // Event offset not committed - will retry on next poll
            }
        }

        consumer.commitSync(); // Commit only if all successful

    } catch (Exception e) {
        logger.error("Consumer error", e);
        // Will reconnect and resume from last committed offset
    }
}

```

Result: Failed events are retried automatically

8. Docker & Containerization

What is Docker?




Simple Analogy: Docker is like **shipping containers for software**: - Just as shipping containers standardize how goods are transported - Docker containers standardize how software runs - Works the same on your laptop, your teammate's computer, and production servers

Without Docker:

Developer 1: "Works on my machine!" (Windows)
 Developer 2: "Doesn't work on mine!" (Mac)
 Server: "Crashes in production!" (Linux)

Problem: Different environments, different results!

With Docker:

Developer 1: Runs Docker container 
 Developer 2: Runs same Docker container 
 Server: Runs same Docker container 

Solution: Same environment everywhere!

Docker Compose

What: Tool to run multiple Docker containers together

File: docker-compose.yml

```
version: '3.8'
```

```
services:
```

```
  # Service 1: Zookeeper (Kafka's coordinator)
```

```
  zookeeper:
```

```
    image: confluentinc/cp-zookeeper:7.5.0
```

```
    ports:
```

```
      - "2181:2181"
```

```
    environment:
```

```
      ZOOKEEPER_CLIENT_PORT: 2181
```

```
    volumes:
```

```
      - zookeeper-data:/var/lib/zookeeper/data
```

```
    healthcheck:
```

```
      test: ["CMD", "nc", "-z", "localhost", "2181"]
```

```
  # Service 2: Kafka (Message broker)
```

```
  kafka:
```

```
    image: confluentinc/cp-kafka:7.5.0
```

```
    depends_on:
```

```
      - zookeeper
```

```
    ports:
```

```
      - "9092:9092"
```

```
    environment:
```

```
      KAFKA_BROKER_ID: 1
```

```
      KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181
```

```
      KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://localhost:9092
```

```
    volumes:
```

```
      - kafka-data:/var/lib/kafka/data
```

```
  # Service 3: PostgreSQL (Database)
```

```
  postgres:
```

```
    image: postgres:16-alpine
```

```
    ports:
```

```
      - "5432:5432"
```

```
    environment:
```

```
      POSTGRES_USER: admin
```

```
      POSTGRES_PASSWORD: admin123
```

```
      POSTGRES_DB: mydb
```

```
    volumes:
```

```
      - postgres-data:/var/lib/postgresql/data
```

```
  # Service 4: MinIO (S3-compatible storage)
```

```
  minio:
```

```
    image: minio/minio:latest
```

```
    command: server /data --console-address ":9001"
```

```
    ports:
```

```
      - "9000:9000" # API
```

```
      - "9001:9001" # Web Console
```

```
    environment:
```

```
      MINIO_ROOT_USER: minioadmin
```

```
      MINIO_ROOT_PASSWORD: minioadmin123
```

```
    volumes:
```

```
      - minio-data:/data
```

```
  # Service 5: pgAdmin (Database UI)
```

```

pgadmin:
  image: dpage/pgadmin4:latest
  ports:
    - "5050:80"
  environment:
    PGADMIN_DEFAULT_EMAIL: admin@admin.com
    PGADMIN_DEFAULT_PASSWORD: admin123

# Service 6: Kafka UI (Kafka management)
kafka-ui:
  image: provectuslabs/kafka-ui:latest
  ports:
    - "8080:8080"
  environment:
    KAFKA_CLUSTERS_0_NAME: local
    KAFKA_CLUSTERS_0_BOOTSTRAPSERVERS: kafka:29092

volumes:
  zookeeper-data:
  kafka-data:
  postgres-data:
  minio-data:
  pgadmin-data:

```

Docker Commands

Start all services:

```

docker-compose up -d
# -d = detached mode (runs in background)

```

Stop all services:

```

docker-compose down

```

Stop and remove all data:

```

docker-compose down -v
# -v = removes volumes (deletes all data)

```

View logs:

```

docker-compose logs kafka
docker-compose logs postgres
docker-compose logs -f # -f = follow (live updates)

```

Check status:

```

docker-compose ps

```

Restart a service:

```

docker-compose restart kafka

```

Docker Networking

Internal Network (docker-compose.yml creates network: kafka-network):

```

Container: kafka-broker
Internal hostname: kafka
Internal IP: 172.18.0.3
Accessible from other containers: kafka:29092

```

Container: postgres-db
Internal hostname: postgres
Internal IP: 172.18.0.4
Accessible from other containers: postgres:5432

Container: minio-s3
Internal hostname: minio
Internal IP: 172.18.0.5
Accessible from other containers: minio:9000

Port Mapping:

Host Machine		Docker Container
localhost:9092	→	kafka:29092
localhost:5432	→	postgres:5432
localhost:9000	→	minio:9000
localhost:8080	→	kafka-ui:8080
localhost:5050	→	pgadmin:80

Volumes & Data Persistence

Without Volumes:

1. Start container
2. Write data
3. Stop container
4. Data is LOST! ❌

With Volumes:

1. Start container (mounts volume)
2. Write data (saved to volume)
3. Stop container
4. Data is PRESERVED! ✅
5. Start container again (data still there)

Our Volumes: - postgres-data → Database files - kafka-data → Kafka logs and events - zookeeper-data → Zookeeper state - minio-data → Uploaded files

9. Deep Dive: Complete Data Flow

Detailed Step-by-Step Execution

Step 1: Application Startup

File: CompleteIntegrationDemo.java:56

```
public static void main(String[] args) {  
    CompleteIntegrationDemo demo = new CompleteIntegrationDemo();  
    demo.initializeComponents();  
    // ...  
}
```

What Happens: 1. JVM starts 2. Loads all classes 3. Initializes logging (Logback) 4. Calls initializeComponents()

Step 2: Initialize S3 Manager

File: S3Manager.java:33

```

public S3Manager(String endpoint, String accessKey, String secretKey) {
    AwsBasicCredentials credentials = AwsBasicCredentials.create(accessKey, secretKey);

    this.s3Client = S3Client.builder()
        .endpointOverride(URI.create("http://localhost:9000"))
        .credentialsProvider(StaticCredentialsProvider.create(credentials))
        .build();
}

```

Network Request:

Java App → MinIO

↳ GET http://localhost:9000/ (health check)
 ◀ Response: 200 OK

Result: S3 client ready to upload files

Step 3: Initialize Database Connection Pool

File: DatabaseManager.java:30

```

HikariConfig config = new HikariConfig();
config.setJdbcUrl("jdbc:postgresql://localhost:5432/mydb");
config.setUsername("admin");
config.setPassword("admin123");
config.setMaximumPoolSize(10);

dataSource = new HikariDataSource(config);

```

Network Request:

Java App → PostgreSQL

↳ TCP connection to localhost:5432
 ◀ Response: Connection accepted
 ↳ AUTH: username=admin, password=***
 ◀ Response: AUTH_OK
 ↳ SELECT 1 (connection test)
 ◀ Response: 1 row

Result: Connection pool with 10 connections ready

Step 4: Initialize Kafka Producer

File: FileEventProducer.java:51

```

Properties props = new Properties();
props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG, StringSerializer.class);
props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG, StringSerializer.class);

producer = new KafkaProducer<>(props);

```

Network Request:

Java App → Kafka

↳ TCP connection to localhost:9092
 ◀ Response: Connection accepted
 ↳ API_VERSIONS request
 ◀ Response: [list of supported APIs]

- ↳ METADATA request (list topics)
 - ◀ Response: {topics: ["file-events"], brokers: [{id:1, host:"kafka"}]}

Result: Producer ready to send messages

Step 5: Initialize Kafka Consumer

File: FileEventConsumer.java:58

```
Properties props = new Properties();
props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
props.put(ConsumerConfig.GROUP_ID_CONFIG, "complete-integration-group");

consumer = new KafkaConsumer<>(props);
consumer.subscribe(Collections.singletonList("file-events"));
```

Network Request:

Java App → Kafka

- ↳ JOIN_GROUP request (group: complete-integration-group)
 - ◀ Response: {memberId: "consumer-1", leader: true}
- ↳ SYNC_GROUP request
 - ◀ Response: {assignment: [file-events-0, file-events-1, file-events-2]}
- ↳ FETCH_OFFSET request (get last read position)
 - ◀ Response: {file-events-0: offset 0, file-events-1: offset 0, file-events-2: offset 0}

Result: Consumer subscribed to topic, ready to receive messages

Step 6: Upload File to S3

File: S3Manager.java:90

```
public String uploadFile(String bucketName, String key, byte[] fileData, String contentType) {
    PutObjectRequest request = PutObjectRequest.builder()
        .bucket("complete-integration-bucket")
        .key("uploads/alice/vacation-photo.jpg")
        .contentType("image/jpeg")
        .contentLength(2500000L)
        .build();

    s3Client.putObject(request, RequestBody.fromBytes(fileData));
}
```

Network Request:

Java App → MinIO

- ↳ PUT http://localhost:9000/complete-integration-bucket/uploads/alice/vacation-photo.jpg
 - Headers:
 - Content-Type: image/jpeg
 - Content-Length: 2500000
 - Authorization: AWS4-HMAC-SHA256 Credential=...
 - Body: [2.5 MB of image data]
 - ◀ Response: 200 OK
 - Headers:
 - ETag: "abc123def456"
 - Location: /complete-integration-bucket/uploads/alice/vacation-photo.jpg

MinIO Internal Processing: 1. Receives HTTP PUT request 2. Validates credentials 3. Writes file to disk: /data/complete-integration-bucket/uploads/alice/vacation-photo.jpg 4. Updates bucket index 5. Calculates ETag (MD5 hash) 6. Returns success response

Result: File stored in MinIO at /data/complete-integration-bucket/uploads/alice/vacation-photo.jpg

Step 7: Save Metadata to PostgreSQL

File: FileRepository.java:48

```
public void insertFileMetadata(FileMetadata metadata) throws SQLException {
    String sql = ""
        INSERT INTO files (id, filename, s3_key, s3_url, file_size,
                           content_type, uploaded_by, uploaded_at)
        VALUES (?, ?, ?, ?, ?, ?, ?, ?)
        "";

    try (Connection conn = databaseManager.getConnection();
         PreparedStatement stmt = conn.prepareStatement(sql)) {

        stmt.setString(1, "a1b2c3d4-e5f6-7g8h-9i0j-k1l2m3n4o5p6");
        stmt.setString(2, "vacation-photo.jpg");
        stmt.setString(3, "uploads/alice/vacation-photo.jpg");
        stmt.setString(4, "http://localhost:9000/complete-integration-
            bucket/uploads/alice/vacation-photo.jpg");
        stmt.setLong(5, 2500000L);
        stmt.setString(6, "image/jpeg");
        stmt.setString(7, "alice");
        stmt.setLong(8, 1729090624308L);

        stmt.executeUpdate();
    }
}
```

Network Request:

Java App → PostgreSQL

↳ TCP packet to localhost:5432

PostgreSQL Wire Protocol:

Message Type: P (Parse)

SQL: INSERT INTO files (id, filename, s3_key, s3_url, file_size,
 content_type, uploaded_by, uploaded_at)
 VALUES (\$1, \$2, \$3, \$4, \$5, \$6, \$7, \$8)

Message Type: B (Bind)

Parameters:

\$1: "a1b2c3d4-e5f6-7g8h-9i0j-k1l2m3n4o5p6"
 \$2: "vacation-photo.jpg"
 \$3: "uploads/alice/vacation-photo.jpg"
 \$4: "http://localhost:9000/complete-integration-bucket/uploads/alice/vacation-photo.jpg"
 \$5: 2500000
 \$6: "image/jpeg"
 \$7: "alice"
 \$8: 1729090624308

Message Type: E (Execute)

← Response:

Message Type: C (CommandComplete)

Tag: INSERT 0 1

PostgreSQL Internal Processing: 1. Parses SQL statement 2. Validates syntax 3. Checks constraints (NOT NULL, PRIMARY KEY) 4. Inserts row into files table 5. Updates indexes (idx_files_uploaded_by,

idx_files_created_at) 6. Writes to Write-Ahead Log (WAL) 7. Returns success

Database State:

```
SELECT * FROM files WHERE id = 'a1b2c3d4-e5f6-7g8h-9i0j-k1l2m3n4o5p6';
```

-- Result:

id	a1b2c3d4-e5f6-7g8h-9i0j-k1l2m3n4o5p6
filename	vacation-photo.jpg
s3_key	uploads/alice/vacation-photo.jpg
s3_url	http://localhost:9000/complete-integration-bucket/uploads/alice/vacation-photo.jpg
file_size	2500000
content_type	image/jpeg
uploaded_by	alice
uploaded_at	1729090624308
created_at	2025-10-16 16:37:04.308
updated_at	2025-10-16 16:37:04.308

Step 8: Publish Event to Kafka

File: FileEventProducer.java:72

```
public void sendUploadedEvent(FileMetadata metadata) throws Exception {
    FileEvent event = new FileEvent("UPLOADED", metadata, System.currentTimeMillis());
    String jsonMessage = objectMapper.writeValueAsString(event);

    ProducerRecord<String, String> record = new ProducerRecord<>(
        "file-events",
        metadata.getId(), // Partition key
        jsonMessage
    );

    producer.send(record).get(); // Wait for acknowledgment
}
```

JSON Message Created:

```
{
  "eventType": "UPLOADED",
  "fileMetadata": {
    "id": "a1b2c3d4-e5f6-7g8h-9i0j-k1l2m3n4o5p6",
    "filename": "vacation-photo.jpg",
    "s3Key": "uploads/alice/vacation-photo.jpg",
    "s3Url": "http://localhost:9000/complete-integration-bucket/uploads/alice/vacation-photo.jpg",
    "fileSize": 2500000,
    "contentType": "image/jpeg",
    "uploadedBy": "alice",
    "uploadedAt": 1729090624308
  },
  "timestamp": 1729090624308
}
```

Network Request:

Java App → Kafka

```
└─► PRODUCE request
    Topic: file-events
    Partition: 1 (calculated from key hash)
    Key: "a1b2c3d4-e5f6-7g8h-9i0j-k1l2m3n4o5p6"
    Value: [JSON message above]
    Compression: none
    Acks: all (wait for all replicas)
```

← Response (after 50ms):
 Status: SUCCESS
 Partition: 1
 Offset: 42
 Timestamp: 1729090624358

Kafka Internal Processing: 1. Receives PRODUCE request 2. Calculates partition from key hash: $\text{hash}(\text{"a1b2c3d4..."}) \% 3 = 1$ 3. Appends message to partition 1 log file 4. Writes to disk 5. Replicates to other brokers (if configured) 6. Returns acknowledgment with offset

Kafka Log File (/var/lib/kafka/data/file-events-1/00000000000000000042.log):

Offset: 42
 Timestamp: 1729090624358
 Key Length: 36
 Key: a1b2c3d4-e5f6-7g8h-9i0j-k1l2m3n4o5p6
 Value Length: 423
 Value: {"eventType":"UPLOADED","fileMetadata":{"...}}

Step 9: Consumer Polls Kafka

File: FileEventConsumer.java:87

```
while (running) {
    ConsumerRecords<String, String> records = consumer.poll(Duration.ofMillis(100));

    for (ConsumerRecord<String, String> record : records) {
        FileEvent event = objectMapper.readValue(record.value(), FileEvent.class);
        eventHandler.accept(event);
    }

    consumer.commitSync();
}
```

Network Request:

Java App → Kafka

↳ FETCH request
 Topics: [file-events]
 Partitions: [0, 1, 2]
 CurrentOffsets: [0, 42, 0] // Partition 1 has new message at offset 42
 MaxWaitMs: 100
 MinBytes: 1
 MaxBytes: 52428800 (50 MB)

← Response (immediately, since message is available):
 Partition: 1
 HighWaterMark: 43
 Messages: [
 {
 Offset: 42
 Key: "a1b2c3d4-e5f6-7g8h-9i0j-k1l2m3n4o5p6"
 Value: "{\\"eventType\\" : \\"UPLOADED\\", \\"fileMetadata\\" : {\\"...\\\"}"
 Timestamp: 1729090624358
 }
]

Java Processing: 1. Receives message from Kafka 2. Deserializes JSON to FileEvent object 3. Calls eventHandler.accept(event) 4. Event handler processes the event

Step 10: Process Event

File: CompleteIntegrationDemo.java:127

```
Consumer<FileEvent> eventHandler = event -> {
    FileMetadata metadata = event.getFileMetadata();

    logger.info("📁 RECEIVED EVENT: {}", event.getEventType());
    logger.info("    • Filename: {}", metadata.getFilename());
    logger.info("    • Size: {}", formatFileSize(metadata.getFileSize()));
    logger.info("    • User: {}", metadata.getUploadedBy());

    // Simulate processing tasks
    logger.info("\n 🔄 Processing tasks:");
    logger.info("    [1] Generating thumbnails...");
    generateThumbnails(metadata);

    logger.info("    [2] Scanning for viruses...");
    scanForViruses(metadata);

    logger.info("    [3] Updating search index...");
    updateSearchIndex(metadata);

    logger.info("    [4] Sending notifications...");
    sendNotification(metadata);

    logger.info("\n ✅ Processing complete!");
    eventsProcessed.incrementAndGet();
};
```

Real-World Processing Examples:

1. Generate Thumbnails:

```
private void generateThumbnails(FileMetadata metadata) {
    if (metadata.getContentType().startsWith("image/")) {
        // Download original image from S3
        byte[] imageData = s3Manager.downloadFile(bucket, metadata.getS3Key());

        // Generate 3 sizes
        byte[] thumb200 = resizeImage(imageData, 200, 200);
        byte[] thumb500 = resizeImage(imageData, 500, 500);
        byte[] thumb1000 = resizeImage(imageData, 1000, 1000);

        // Upload thumbnails back to S3
        s3Manager.uploadFile(bucket, metadata.getS3Key() + "_thumb_200", thumb200, "image/jpeg");
        s3Manager.uploadFile(bucket, metadata.getS3Key() + "_thumb_500", thumb500, "image/jpeg");
        s3Manager.uploadFile(bucket, metadata.getS3Key() + "_thumb_1000", thumb1000,
            "image/jpeg");
    }
}
```

2. Scan for Viruses:

```
private void scanForViruses(FileMetadata metadata) {
    // Download file from S3
    byte[] fileData = s3Manager.downloadFile(bucket, metadata.getS3Key());

    // Call virus scanning API (e.g., ClamAV)
    boolean isClean = virusScanner.scan(fileData);

    if (!isClean) {
        // Quarantine file
        s3Manager.deleteFile(bucket, metadata.getS3Key());
        fileRepository.deleteFileMetadata(metadata.getId());

        // Notify user
    }
}
```

```

        emailService.send(metadata.getUploadedBy(),
            "Your file was removed due to security concerns");
    }
}

```

3. Update Search Index:

```

private void updateSearchIndex(FileMetadata metadata) {
    // Create search document
    SearchDocument doc = new SearchDocument();
    doc.setId(metadata.getId());
    doc.setTitle(metadata.getFilename());
    doc.setContent(extractText(metadata)); // OCR for images, parse PDF, etc.
    doc.setUser(metadata.getUploadedBy());
    doc.setTimestamp(metadata.getUploadedAt());

    // Index in Elasticsearch
    elasticsearchClient.index("files", doc);
}

```

4. Send Notification:

```

private void sendNotification(FileMetadata metadata) {
    // Send email
    emailService.send(
        metadata.getUploadedBy(),
        "File uploaded successfully",
        "Your file '" + metadata.getFilename() + "' is ready!"
    );

    // Push notification
    pushService.send(
        metadata.getUploadedBy(),
        "File uploaded",
        metadata.getFilename()
    );

    // Update user's notification feed
    notificationRepository.create(
        metadata.getUploadedBy(),
        "FILE_UPLOADED",
        metadata.getId()
    );
}

```

Step 11: Commit Offset

File: FileEventConsumer.java:102

```
consumer.commitSync();
```

Network Request:

Java App → Kafka

```

└─► OFFSET_COMMIT request
    Group: complete-integration-group
    Offsets: {
        file-events-0: 0,
        file-events-1: 43, // We processed message at offset 42
        file-events-2: 0
    }

```

← Response:
Status: SUCCESS

What This Means: - If consumer crashes and restarts, it will resume from offset 43 - Message at offset 42 won't be processed again - Ensures "exactly-once" processing (with proper configuration)

Complete Timeline Summary

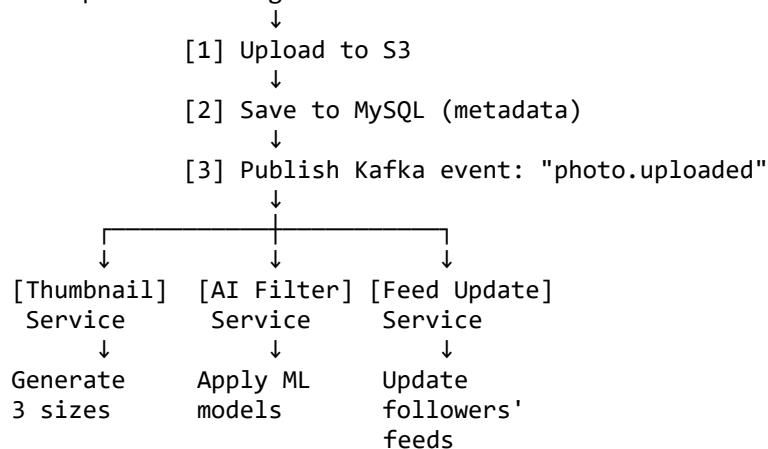
Time	Component	Action	Duration
0.000s	User	Clicks "Upload"	-
0.001s	Browser	Sends HTTP POST to Java app	-
0.002s	Java App	Receives file (10 MB)	-
0.003s	S3Manager	Starts upload to MinIO	-
0.500s	MinIO	File uploaded	497ms
0.501s	FileRepository	Starts DB insert	-
0.520s	PostgreSQL	Row inserted	19ms
0.521s	FileEventProducer	Creates JSON event	-
0.522s	Kafka	Event stored	1ms
0.523s	Java App	Returns HTTP 200 to user	-
0.524s	User	Sees "Upload successful!"	Total: 524ms
1.000s	Consumer	Polls Kafka	-
1.001s	Kafka	Returns event	1ms
1.002s	Event Handler	Starts processing	-
2.000s	Thumbnail Gen	Generates 3 thumbnails	998ms
10.000s	Virus Scanner	Scans file	8000ms
12.000s	Search Indexer	Updates Elasticsearch	2000ms
13.000s	Notification	Sends email & push	1000ms
13.001s	Consumer	Commits offset	-
TOTAL	User perceived time	524ms	
TOTAL	Background processing	12.5 seconds (asynchronous)	

10. Real-World Applications

Case Study 1: Instagram Photo Upload

Architecture:

User uploads photo → Instagram API



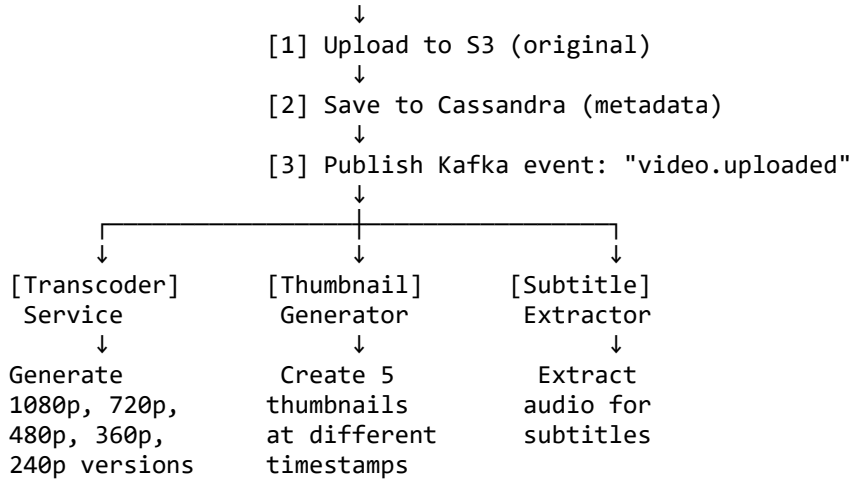
Our Implementation Maps To: - `S3Manager.uploadFile()` → Instagram's photo storage - `FileRepository.insertFileMetadata()` → Instagram's post metadata - `FileEventProducer.sendUploadedEvent()` → Instagram's event system - `FileEventConsumer` → Instagram's background processors

Scale: - 95 million photos uploaded per day - 1,100 photos per second - Each photo generates ~20 events - 22,000 events per second processed

Case Study 2: Netflix Video Upload

Architecture:

Content provider uploads video → Netflix API



Transcoding Pipeline:

// Similar to our event handler

```

Consumer<VideoEvent> transcodingHandler = event -> {
    VideoMetadata video = event.getVideoMetadata();

    // Download original from S3
    byte[] originalVideo = s3Manager.downloadFile(bucket, video.getS3Key());

    // Transcode to multiple qualities
    byte[] video1080p = transcode(originalVideo, "1080p");
    byte[] video720p = transcode(originalVideo, "720p");
    byte[] video480p = transcode(originalVideo, "480p");
    byte[] video360p = transcode(originalVideo, "360p");
    byte[] video240p = transcode(originalVideo, "240p");

    // Upload all versions back to S3
    s3Manager.uploadFile(bucket, video.getS3Key() + "_1080p", video1080p, "video/mp4");
    s3Manager.uploadFile(bucket, video.getS3Key() + "_720p", video720p, "video/mp4");
    // ... and so on

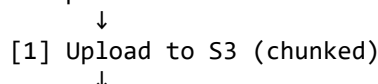
    // Publish event: "video.transcoded"
    eventProducer.sendTranscodedEvent(video);
};
  
```

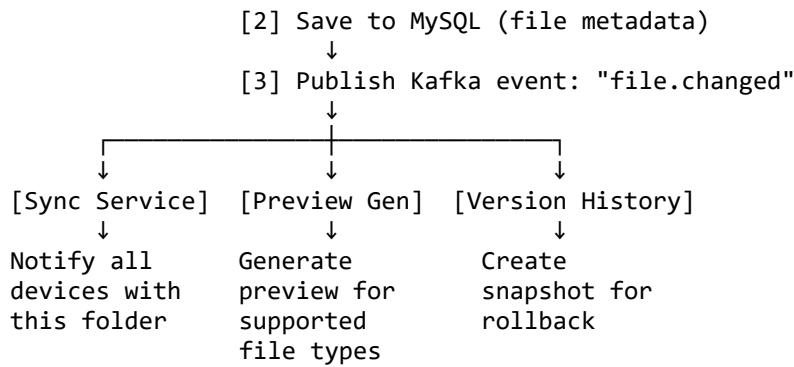
Scale: - Netflix processes thousands of hours of new content per month - Each video generates 5 quality versions + multiple thumbnails - Transcoding can take hours for a 2-hour movie - All done asynchronously while content provider sees immediate confirmation

Case Study 3: Dropbox File Sync

Architecture:

User uploads file to Dropbox → Dropbox API





File Chunking (How Dropbox uploads large files):

```

public void uploadLargeFile(String filename, byte[] fileData) {
    int CHUNK_SIZE = 4 * 1024 * 1024; // 4 MB chunks
    int numChunks = (int) Math.ceil((double) fileData.length / CHUNK_SIZE);

    String uploadSessionId = UUID.randomUUID().toString();

    for (int i = 0; i < numChunks; i++) {
        int start = i * CHUNK_SIZE;
        int end = Math.min(start + CHUNK_SIZE, fileData.length);
        byte[] chunk = Arrays.copyOfRange(fileData, start, end);

        // Upload chunk
        s3Manager.uploadFilePart(
            bucket,
            uploadSessionId,
            i, // Part number
            chunk
        );

        logger.info("Uploaded chunk {}/{}", i + 1, numChunks);
    }

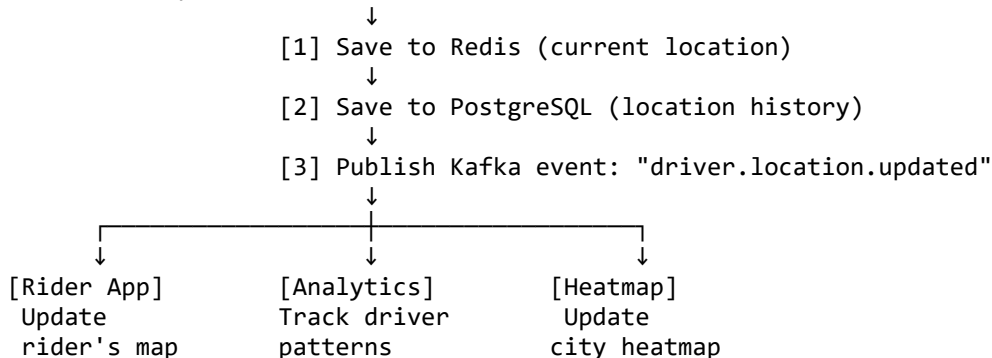
    // Complete multipart upload
    String s3Key = s3Manager.completeMultipartUpload(uploadSessionId, filename);

    // Save metadata and publish event
    saveMetadataAndPublishEvent(s3Key, filename, fileData.length);
}
  
```

Case Study 4: Uber Ride Tracking

Architecture (Simplified):

Driver's phone sends GPS update → Uber API



Real-time Location Updates:

```

Consumer<LocationEvent> locationHandler = event -> {
    DriverLocation location = event.getLocation();

    // Update Redis (fast, in-memory)
    redisClient.set("driver:" + location.getDriverId(), location);

    // Save to PostgreSQL (persistent history)
    locationRepository.insertLocation(location);

    // Find nearby riders looking for rides
    List<Rider> nearbyRiders = findRidersNearLocation(location);

    // Notify each rider
    for (Rider rider : nearbyRiders) {
        pushNotification(rider, "Driver nearby!");
    }

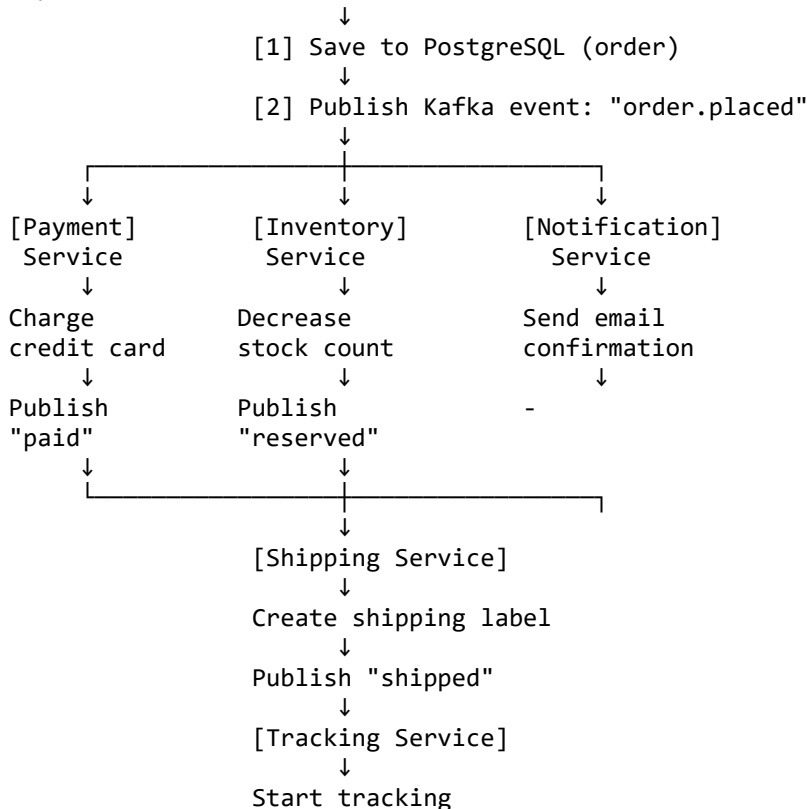
    // Update city heatmap
    heatmapService.updateCell(location.getLat(), location.getLon());
};

```

Case Study 5: E-commerce Order Processing

Architecture:

User places order → E-commerce API



Order Processing Flow:

```

// Order placement
public void placeOrder(Order order) {
    // Save order to database
    orderRepository.insertOrder(order);

    // Publish event
    kafkaProducer.send("order-events", new OrderEvent("PLACED", order));
}

```

```
// Payment processor
Consumer<OrderEvent> paymentHandler = event -> {
    if (event.getType().equals("PLACED")) {
        Order order = event.getOrder();

        // Charge credit card
        boolean success = paymentGateway.charge(
            order.getPaymentMethod(),
            order.getTotalAmount()
        );

        if (success) {
            // Update order status
            orderRepository.updateStatus(order.getId(), "PAID");

            // Publish payment success event
            kafkaProducer.send("order-events",
                new OrderEvent("PAID", order));
        } else {
            // Handle payment failure
            orderRepository.updateStatus(order.getId(), "PAYMENT_FAILED");
            kafkaProducer.send("order-events",
                new OrderEvent("PAYMENT_FAILED", order));
        }
    }
};

// Inventory manager
Consumer<OrderEvent> inventoryHandler = event -> {
    if (event.getType().equals("PAID")) {
        Order order = event.getOrder();

        // Decrease stock for each item
        for (OrderItem item : order.getItems()) {
            inventoryRepository.decreaseStock(
                item.getProductId(),
                item.getQuantity()
            );
        }

        // Publish inventory reserved event
        kafkaProducer.send("order-events",
            new OrderEvent("INVENTORY_RESERVED", order));
    }
};

// Shipping service
Consumer<OrderEvent> shippingHandler = event -> {
    if (event.getType().equals("INVENTORY_RESERVED")) {
        Order order = event.getOrder();

        // Create shipping label
        ShippingLabel label = shippingService.createLabel(order);

        // Update order with tracking number
        orderRepository.updateTrackingNumber(
            order.getId(),
            label.getTrackingNumber()
        );

        // Publish shipped event
        kafkaProducer.send("order-events",
            new OrderEvent("SHIPPED", order));
    }
};
```

```
// Notification service
Consumer<OrderEvent> notificationHandler = event -> {
    Order order = event.getOrder();

    switch (event.getType()) {
        case "PLACED":
            emailService.send(order.getCustomerEmail(),
                "Order received",
                "We've received your order #" + order.getId());
            break;

        case "PAID":
            emailService.send(order.getCustomerEmail(),
                "Payment confirmed",
                "Your payment has been processed");
            break;

        case "SHIPPED":
            emailService.send(order.getCustomerEmail(),
                "Order shipped",
                "Track your order: " + order.getTrackingNumber());
            break;

        case "PAYMENT_FAILED":
            emailService.send(order.getCustomerEmail(),
                "Payment failed",
                "Please update your payment method");
            break;
    }
};
```

Why Event-Driven Architecture Wins Here: - **Decoupling:** Payment failure doesn't crash inventory system - **Scalability:** Each service scales independently - **Resilience:** If email service is down, order still processes - **Auditability:** Complete event log of order lifecycle - **Flexibility:** Easy to add new services (fraud detection, recommendations, etc.)

11. Key Concepts for Beginners

1. Event-Driven Architecture

Traditional (Synchronous):

```
// Caller waits for each step to complete
void uploadFile(File file) {
    s3.upload(file);           // Wait 3 seconds
    database.save(metadata);   // Wait 0.5 seconds
    thumbnail.generate(file);  // Wait 10 seconds
    virus.scan(file);          // Wait 15 seconds
    email.send(notification);   // Wait 2 seconds
    // Total wait: 30.5 seconds 🐢
}
```

Event-Driven (Asynchronous):

```
// Caller returns immediately
void uploadFile(File file) {
    s3.upload(file);
    database.save(metadata);
    kafka.publish("file.uploaded", metadata);
    // Returns immediately! ⚡ (3.5 seconds)
}
```



```
// Background services process independently
void onFileUploaded(FileEvent event) {
    thumbnail.generate(event.file); // Service 1
}

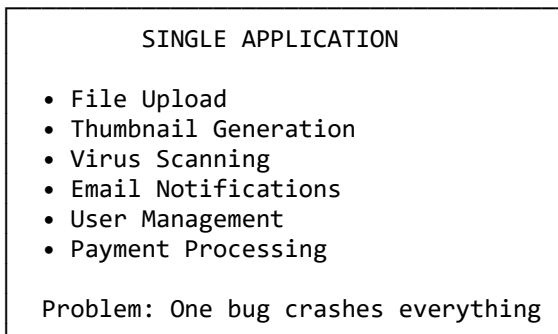
void onFileUploaded(FileEvent event) {
    virus.scan(event.file); // Service 2
}

void onFileUploaded(FileEvent event) {
    email.send(notification); // Service 3
}
```

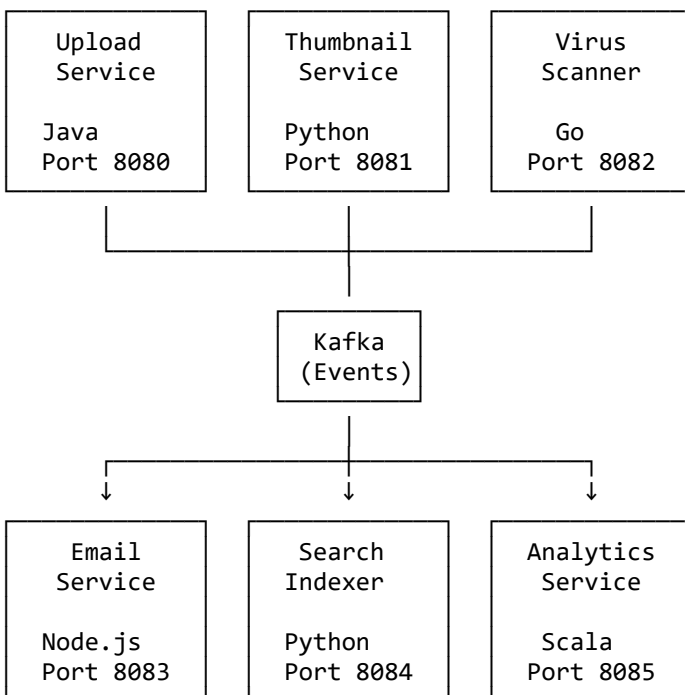
Benefits: - ☒ Fast response to user - ☒ Services don't block each other - ☒ Easy to add new services - ☒ Handles failures gracefully

2. Microservices Architecture

Monolith (All in one):



Microservices (Separate services):



Benefits: - ☒ Independent deployment - ☒ Technology flexibility - ☒ Team autonomy - ☒ Fault isolation - ☒ Independent scaling

3. CAP Theorem

CAP Theorem: In a distributed system, you can only guarantee 2 out of 3:

1. **Consistency:** All nodes see the same data at the same time
2. **Availability:** Every request receives a response
3. **Partition Tolerance:** System continues despite network failures

Examples:

PostgreSQL (CP - Consistency + Partition Tolerance):

Scenario: Network partition between database replicas

Choice: Reject writes to maintain consistency

Result: Some requests fail, but data stays consistent

Use case: Banking (consistency critical)

Cassandra (AP - Availability + Partition Tolerance):

Scenario: Network partition between database nodes

Choice: Accept writes on all nodes

Result: All requests succeed, but data may be inconsistent temporarily

Use case: Social media (availability critical)

Our Architecture: - **PostgreSQL**: CP (consistency for metadata) - **S3/MinIO**: AP (availability for file storage) - **Kafka**: AP (availability for event delivery)

4. ACID vs BASE

ACID (Traditional databases): - **Atomicity**: All or nothing - **Consistency**: Data always valid - **Isolation**: Transactions don't interfere - **Durability**: Committed data persists

Example:

```
// Bank transfer - ACID
beginTransaction();
try {
    account1.subtract(100); // Must succeed
    account2.add(100);      // Must succeed
    commit();               // Both succeed or both fail
} catch (Exception e) {
    rollback();             // Undo everything
}
```

BASE (Modern distributed systems): - **Basically Available**: System is always available - **Soft state**: State may change without input - **Eventual consistency**: Data becomes consistent eventually

Example:

```
// File upload - BASE
s3.upload(file); // Available immediately
kafka.publish("file.uploaded"); // Eventually processed
// Thumbnails generated eventually
// Search index updated eventually
// User notified eventually
```

5. Connection Pooling

Without Pooling (Slow):

```
// Every request creates a new connection
void handleRequest() {
    Connection conn = DriverManager.getConnection(url, user, pass);
    // Connection creation takes 100ms!

    // Do work (10ms)
    Statement stmt = conn.createStatement();
    ResultSet rs = stmt.executeQuery("SELECT * FROM files");

    conn.close();
    // Total: 110ms per request
}
```





With Pooling (Fast):

```
// Connections created once at startup
HikariCP pool = new HikariCP(10); // 10 connections ready

void handleRequest() {
    Connection conn = pool.getConnection(); // Instant! (0.1ms)

    // Do work (10ms)
    Statement stmt = conn.createStatement();
    ResultSet rs = stmt.executeQuery("SELECT * FROM files");

    conn.close(); // Returns to pool, doesn't actually close
    // Total: 10ms per request (11x faster!)
}
```

Benefits: -  10-100x faster -  Less resource usage -  Automatic connection management - 
Connection health checks

6. Serialization & Deserialization

Serialization: Converting objects to bytes/string **Deserialization:** Converting bytes/string back to objects

Example with JSON:

```
// Java Object
FileMetadata metadata = new FileMetadata(
    "123",
    "photo.jpg",
    "uploads/photo.jpg",
    "http://s3.com/photo.jpg",
    2500000L,
    "image/jpeg",
    "alice",
    1729090624308L
);

// Serialization (Object → JSON String)
ObjectMapper mapper = new ObjectMapper();
String json = mapper.writeValueAsString(metadata);
// Result: {"id":"123","filename":"photo.jpg",...}

// Send JSON over network to Kafka

// Deserialization (JSON String → Object)
FileMetadata received = mapper.readValue(json, FileMetadata.class);
// Result: FileMetadata object with same values
```

Why JSON? -  Human-readable -  Language-independent -  Widely supported -  Easy to debug

Alternatives: - **Protocol Buffers:** Faster, smaller, but binary - **Avro:** Schema evolution support - **MessagePack:** Compact binary format

7. Idempotency

Idempotent: Same operation can be repeated without different outcome

Non-Idempotent (Bad):

```
void processPayment(Order order) {  
    // If this runs twice, charges customer twice! ❌  
    creditCard.charge(order.getTotalAmount());  
}
```

Idempotent (Good):

```
void processPayment(Order order) {  
    // Check if already processed  
    if (paymentRepository.exists(order.getId())) {  
        return; // Already processed, skip  
    }  
  
    // Process payment  
    creditCard.charge(order.getTotalAmount());  
  
    // Mark as processed  
    paymentRepository.markProcessed(order.getId());  
}
```

In Kafka:

```
// Enable idempotence  
props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, "true");  
  
// Kafka prevents duplicate sends automatically  
// If network fails and retries, Kafka detects duplicate
```

8. Eventual Consistency

Strong Consistency (Traditional):

```
User uploads file  
↓  
Write to database (wait for replication to all nodes)  
↓  
Read from database (always sees latest data)
```

Eventual Consistency (Distributed):

```
User uploads file  
↓  
Write to database (write to one node, return immediately)  
↓  
Background: Replicate to other nodes  
↓  
Eventually: All nodes have the same data
```

Example in Our System:

```
Time 0:00 - User uploads file  
Time 0:01 - File saved to S3, metadata saved to PostgreSQL
```

Time 0:02 - Kafka event published
Time 0:03 - Consumer starts processing
Time 0:05 - Thumbnails generated
Time 0:10 - Search index updated
Time 0:15 - All systems consistent ✅

During 0:01-0:15: - File is uploaded ✅ - Metadata in database ✅ - But thumbnails not ready yet ⌚ -
And search doesn't show it yet ⌚

Eventually (at 0:15): - Everything is consistent ✅

9. Backpressure

Problem: Producer sends messages faster than consumer can process

Without Backpressure (Bad):

Producer: 1000 messages/second
Consumer: 100 messages/second
Result: Queue grows infinitely, memory overflow! ✖

With Backpressure (Good):

Consumer tells producer: "Slow down, I can only handle 100/sec"
Producer throttles to match consumer's capacity
Result: Stable system ✅

In Kafka:

```
// Consumer automatically handles backpressure
consumer.poll(Duration.ofMillis(100)); // Fetch as much as can process
// Kafka won't send more until consumer commits offset
```

10. Schema Evolution

Problem: How to update data structure without breaking existing code?

Example:

```
// Version 1
class FileMetadata {
    String id;
    String filename;
}

// Deployed to production, messages sent

// Version 2 (need to add new field)
class FileMetadata {
    String id;
    String filename;
    String contentType; // NEW FIELD
}

// Problem: Old messages don't have contentType!
```

Solution 1: Optional Fields

```
class FileMetadata {
    String id;
```

```
String filename;  
String contentType = "application/octet-stream"; // Default value  
}
```

Solution 2: Schema Registry (Avro, Protobuf):

```
// Schema Registry tracks versions  
// V1: {id, filename}  
// V2: {id, filename, contentType}  
// Consumer knows how to read both versions
```

12. Troubleshooting & Common Issues

Issue 1: Kafka Connection Refused

Error:

```
org.apache.kafka.common.errors.TimeoutException:  
Failed to update metadata after 60000 ms.
```

Cause: Kafka not running or wrong host/port

Solution:

```
# Check if Kafka is running  
docker-compose ps  
  
# Should see:  
# kafka-broker    Up (healthy)  
  
# If not running:  
docker-compose up -d kafka  
  
# Check logs:  
docker-compose logs kafka
```

Issue 2: PostgreSQL Connection Failed

Error:

```
org.postgresql.util.PSQLException:  
Connection refused. Check that the hostname and port are correct.
```

Cause: PostgreSQL not running or wrong credentials

Solution:

```
# Check if PostgreSQL is running  
docker-compose ps postgres  
  
# Test connection:  
docker-compose exec postgres psql -U admin -d mydb  
  
# If password error, check docker-compose.yml:  
postgres:  
  environment:  
    POSTGRES_USER: admin          # Must match code  
    POSTGRES_PASSWORD: admin123  # Must match code
```

Issue 3: MinIO Access Denied

Error:

software.amazon.awssdk.services.s3.model.S3Exception:
Access Denied (Service: S3, Status Code: 403)

Cause: Wrong access key/secret key

Solution:

```
// Check credentials in code
S3Manager s3Manager = new S3Manager(
    "http://localhost:9000",
    "minioadmin",           // Must match docker-compose.yml
    "minioadmin123"        // Must match docker-compose.yml
);

# Check docker-compose.yml
minio:
  environment:
    MINIO_ROOT_USER: minioadmin
    MINIO_ROOT_PASSWORD: minioadmin123
```

Issue 4: Out of Memory

Error:

java.lang.OutOfMemoryError: Java heap space

Cause: Uploading too large files or too many at once

Solution 1: Increase JVM memory

```
# Add to run command
java -Xmx2g -Xms512m -jar app.jar
# -Xmx2g = max 2 GB
# -Xms512m = start with 512 MB
```

Solution 2: Stream large files

```
// BAD: Loads entire file into memory
byte[] fileData = Files.readAllBytes(path); // 1 GB file = 1 GB memory!

// GOOD: Stream file in chunks
try (InputStream is = Files.newInputStream(path)) {
    s3Client.putObject(request, RequestBody.fromInputStream(is, fileSize));
}
```

Issue 5: Kafka Lag (Consumer Too Slow)

Symptom: Events processed with significant delay

Diagnosis:

```
# Check consumer lag
docker-compose exec kafka kafka-consumer-groups \
  --bootstrap-server localhost:9092 \
  --describe \
```

```
--group complete-integration-group
```

Output shows:

```
# GROUP          TOPIC      PARTITION  CURRENT-OFFSET  LAG
# complete-integration-group file-events 0          1000            500
# complete-integration-group file-events 1          1000            500
# complete-integration-group file-events 2          1000            500
# Total lag: 1500 messages behind!
```

Solution 1: Add more consumers

```
// Run multiple instances
// They'll automatically share the work
// Max consumers = number of partitions (3 in our case)
```

```
// Terminal 1
java -jar app.jar
```

```
// Terminal 2
java -jar app.jar
```

```
// Terminal 3
java -jar app.jar
```

Solution 2: Optimize processing

```
// BAD: Slow processing
Consumer<FileEvent> handler = event -> {
    // Synchronous processing (slow)
    generateThumbnails(event); // 10 seconds
    scanVirus(event);         // 15 seconds
    sendEmail(event);          // 2 seconds
    // Total: 27 seconds per event!
};

// GOOD: Parallel processing
Consumer<FileEvent> handler = event -> {
    // Async processing (fast)
    CompletableFuture.runAsync(() -> generateThumbnails(event));
    CompletableFuture.runAsync(() -> scanVirus(event));
    CompletableFuture.runAsync(() -> sendEmail(event));
    // Total: Max 15 seconds per event (3 tasks in parallel)
};
```

Issue 6: Database Deadlock

Error:

```
org.postgresql.util.PSQLException:
ERROR: deadlock detected
```

Cause: Two transactions waiting for each other

Example:

```
Transaction 1: Lock file A, waiting for file B
Transaction 2: Lock file B, waiting for file A
Result: Deadlock! 💀
```

Solution: Consistent lock ordering


```
// BAD: Random order
void transferFiles(String fileId1, String fileId2) {
    lockFile(fileId1);
    lockFile(fileId2);
    // ...
}

// GOOD: Always lock in same order
void transferFiles(String fileId1, String fileId2) {
    String first = fileId1.compareTo(fileId2) < 0 ? fileId1 : fileId2;
    String second = fileId1.compareTo(fileId2) < 0 ? fileId2 : fileId1;

    lockFile(first);
    lockFile(second);
    // ...
}
```

Issue 7: S3 Upload Timeout

Error:

software.amazon.awssdk.core.exception.SdkClientException:
Unable to execute HTTP request: Read timed out

Cause: Large file upload taking too long

Solution 1: Increase timeout

```
S3Client s3Client = S3Client.builder()
    .overrideConfiguration(config -> config
        .apiCallTimeout(Duration.ofMinutes(10)) // 10 min timeout
        .apiCallAttemptTimeout(Duration.ofMinutes(5)))
    .build();
```

Solution 2: Multipart upload for large files

```
public void uploadLargeFile(String bucket, String key, File file) {
    CreateMultipartUploadRequest createRequest = CreateMultipartUploadRequest.builder()
        .bucket(bucket)
        .key(key)
        .build();

    String uploadId = s3Client.createMultipartUpload(createRequest).uploadId();

    // Upload in 5 MB chunks
    int partSize = 5 * 1024 * 1024;
    int partNumber = 1;
    List<CompletedPart> completedParts = new ArrayList<>();

    try (FileInputStream fis = new FileInputStream(file)) {
        byte[] buffer = new byte[partSize];
        int bytesRead;

        while ((bytesRead = fis.read(buffer)) > 0) {
            UploadPartRequest uploadRequest = UploadPartRequest.builder()
                .bucket(bucket)
                .key(key)
                .uploadId(uploadId)
                .partNumber(partNumber)
                .build();

            UploadPartResponse response = s3Client.uploadPart(uploadRequest,
                RequestBody.fromBytes(ByteArrays.copyOf(buffer, bytesRead)));
        }
    }
}
```

```

        completedParts.add(CompletedPart.builder()
            .partNumber(partNumber)
            .eTag(response.eTag())
            .build());

        partNumber++;
    }
}

// Complete upload
CompleteMultipartUploadRequest completeRequest = CompleteMultipartUploadRequest.builder()
    .bucket(bucket)
    .key(key)
    .uploadId(uploadId)
    .multipartUpload(CompletedMultipartUpload.builder()
        .parts(completedParts)
        .build())
    .build();

s3Client.completeMultipartUpload(completeRequest);
}

```

Issue 8: Consumer Rebalancing Loop

Symptom: Consumer constantly disconnecting and reconnecting

Logs:

```

[INFO] Revoking previously assigned partitions [file-events-0, file-events-1]
[INFO] Setting newly assigned partitions [file-events-0, file-events-1]
[INFO] Revoking previously assigned partitions [file-events-0, file-events-1]
[INFO] Setting newly assigned partitions [file-events-0, file-events-1]
// Repeats forever...

```

Cause: Consumer processing takes longer than session timeout

Solution: Increase timeouts

```

Properties props = new Properties();
props.put(ConsumerConfig.SESSION_TIMEOUT_MS_CONFIG, "30000"); // 30 seconds
props.put(ConsumerConfig.MAX_POLL_INTERVAL_MS_CONFIG, "300000"); // 5 minutes
props.put(ConsumerConfig.MAX_POLL_RECORDS_CONFIG, "10"); // Process fewer records

consumer = new KafkaConsumer<>(props);

```

Debugging Checklist

When something doesn't work:

- ☐ 1. Check all services are running
 docker-compose ps
 (All should show "Up" or "Up (healthy)")
- ☐ 2. Check service logs
 docker-compose logs kafka
 docker-compose logs postgres
 docker-compose logs minio
- ☐ 3. Check network connectivity
 # From your Java app, can you reach services?

```
telnet localhost 9092 # Kafka
telnet localhost 5432 # PostgreSQL
telnet localhost 9000 # MinIO
```

- ☐ 4. Check credentials
 - # Are usernames/passwords correct?
 - # Check docker-compose.yml matches your code
 - ☐ 5. Check disk space
 - df -h
 - # Ensure you have enough space for data
 - ☐ 6. Check Java version
 - java -version
 - # Should be Java 11 or higher
 - ☐ 7. Check Maven dependencies
 - mvn dependency:tree
 - # Ensure all dependencies are downloaded
 - ☐ 8. Clean and rebuild
 - mvn clean compile
 - docker-compose down -v
 - docker-compose up -d
-

Conclusion

Congratulations! You now understand:

1. ☒ **Apache Kafka** - Message streaming and event processing
2. ☒ **PostgreSQL** - Relational database and metadata storage
3. ☒ **MinIO (S3)** - Blob storage for large files
4. ☒ **Java Application** - Orchestrating all components
5. ☒ **Docker** - Containerization and deployment
6. ☒ **Event-Driven Architecture** - Building scalable systems
7. ☒ **Real-World Applications** - How Netflix, Instagram, Dropbox work
8. ☒ **Troubleshooting** - Solving common issues

Next Steps

1. **Extend the project:**
 - Add user authentication
 - Implement file sharing
 - Add thumbnail generation
 - Create a REST API
 - Build a web frontend
2. **Learn more:**
 - Kubernetes (container orchestration)
 - Redis (caching)
 - Elasticsearch (search)
 - GraphQL (API design)
 - React (frontend)
3. **Deploy to production:**
 - Use AWS S3 instead of MinIO
 - Use AWS MSK instead of local Kafka
 - Use AWS RDS instead of local PostgreSQL
 - Add monitoring (Prometheus, Grafana)
 - Add logging (ELK stack)

Resources

- **Kafka:** <https://kafka.apache.org/documentation/>
 - **PostgreSQL:** <https://www.postgresql.org/docs/>
 - **AWS S3:** <https://docs.aws.amazon.com/s3/>
 - **Docker:** <https://docs.docker.com/>
 - **Java:** <https://docs.oracle.com/en/java/>
-

Thank you for learning with this guide! 🎉

If you have questions or find issues, please refer to the README.md and other documentation files in the project.